



Utilizing Apache Pulsar, Apache  
NiFi and MiNiFi for EdgeAI IoT  
at Scale

**Tim Spann** | Developer Advocate



**Tim Spann**  
Developer Advocate

DZone Zone Leader and Big Data  
MVB Data DJay

- <https://www.datainmotion.dev/>
- <https://github.com/tspannhw/SpeakerProfile>
- <https://dev.to/tspannhw>
- <https://sessionize.com/tspann/>





Founded by the original developers of Apache Pulsar and Apache BookKeeper, StreamNative builds a cloud-native event streaming platform that enables enterprises to easily access data as real-time event streams.

# Apache Pulsar



**Apache**  **PULSAR** is an open source, cloud-native distributed messaging and streaming platform.

# What are the Benefits of Pulsar?



Multi-Tenancy

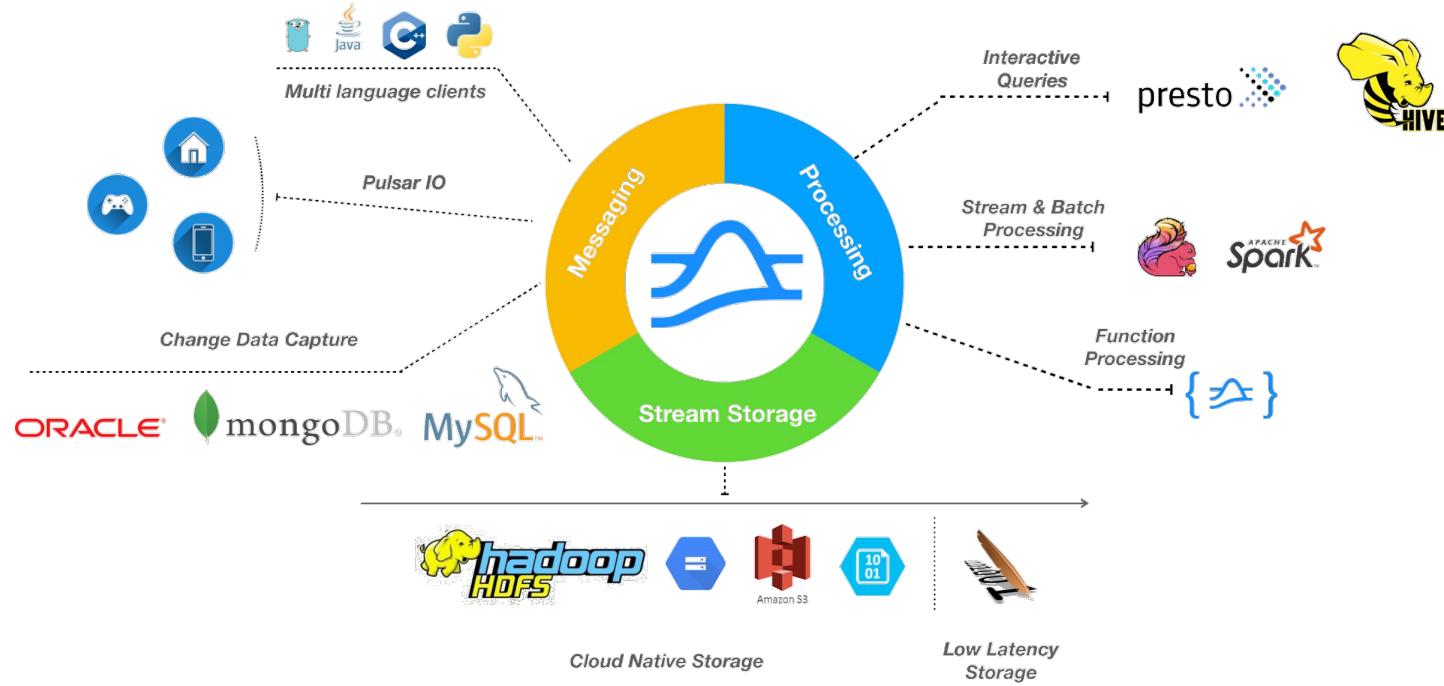
Scalability

Geo-Replication

Unified Messaging  
Model

Data Durability

# Apache Pulsar



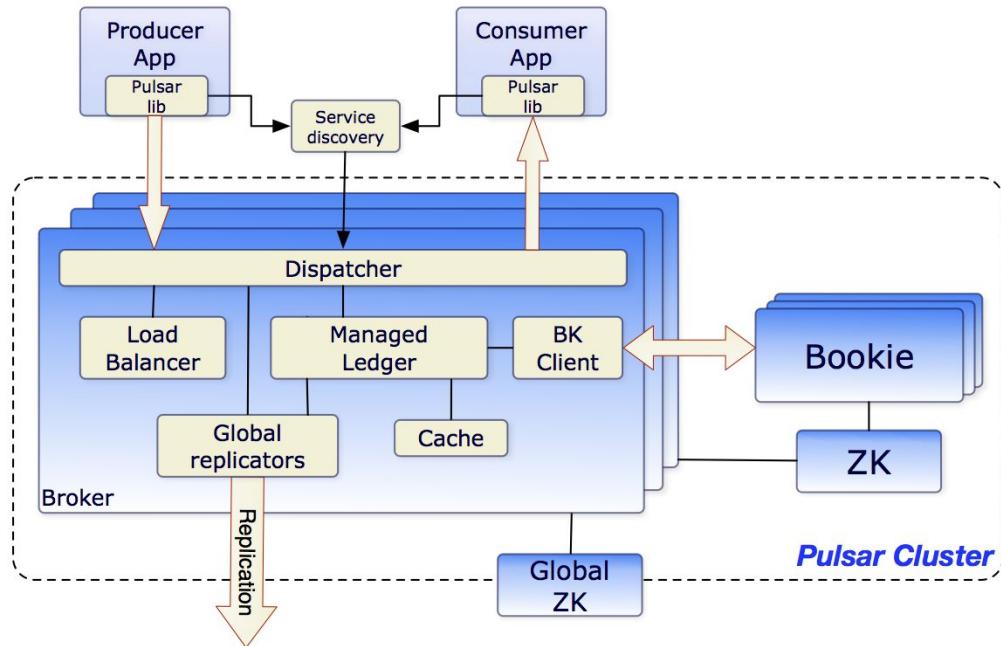
# A Unified Messaging Platform



# Apache Pulsar Overview

## Enable Geo-Replicated Messaging

- Pub-Sub
- Geo-Replication
- Pulsar Functions
- Horizontal Scalability
- Multi-tenancy
- Tiered Persistent Storage
- Pulsar Connectors
- REST API
- CLI
- Many clients available
- Four Different Subscription Types
- Multi-Protocol Support
  - MQTT
  - AMQP
  - JMS
  - Kafka
  - ...



# What is the Pulsar Ecosystem?

- **Functions and Connectors**
  - Functions: Lightweight stream processing
  - Connectors: Part of “Pulsar IO”, includes “Source” and “Sink” APIs
    - Files, Databases, Data tools, Cloud Services, etc
- **Protocol Handlers**
  - Allows Pulsar to handle additional protocols by an extendable API running in the broker
    - AoP (AMQP), KoP (Kafka), MoP (MQTT)

# What is the Pulsar Ecosystem? (cont'd)

- **Processing Engines**
  - Supports modern processing engines
    - Flink and Spark, as well as Pulsar SQL (Presto/Trino)
- **Offloaders**
  - Allows data to be offloaded to cloud storage and used with existing Pulsar APIs
    - S3, GCP Cloud Storage, HDFS, File (NFS), Azure Blob Storage (in Pulsar 2.7.0)

# Pulsar Functions

Provides a simple API to:

- Receive a message (consume)
- Process the message using your own code
- Send a message (produce)

Takes care of the boilerplate code so there is no need to create producers and consumers.

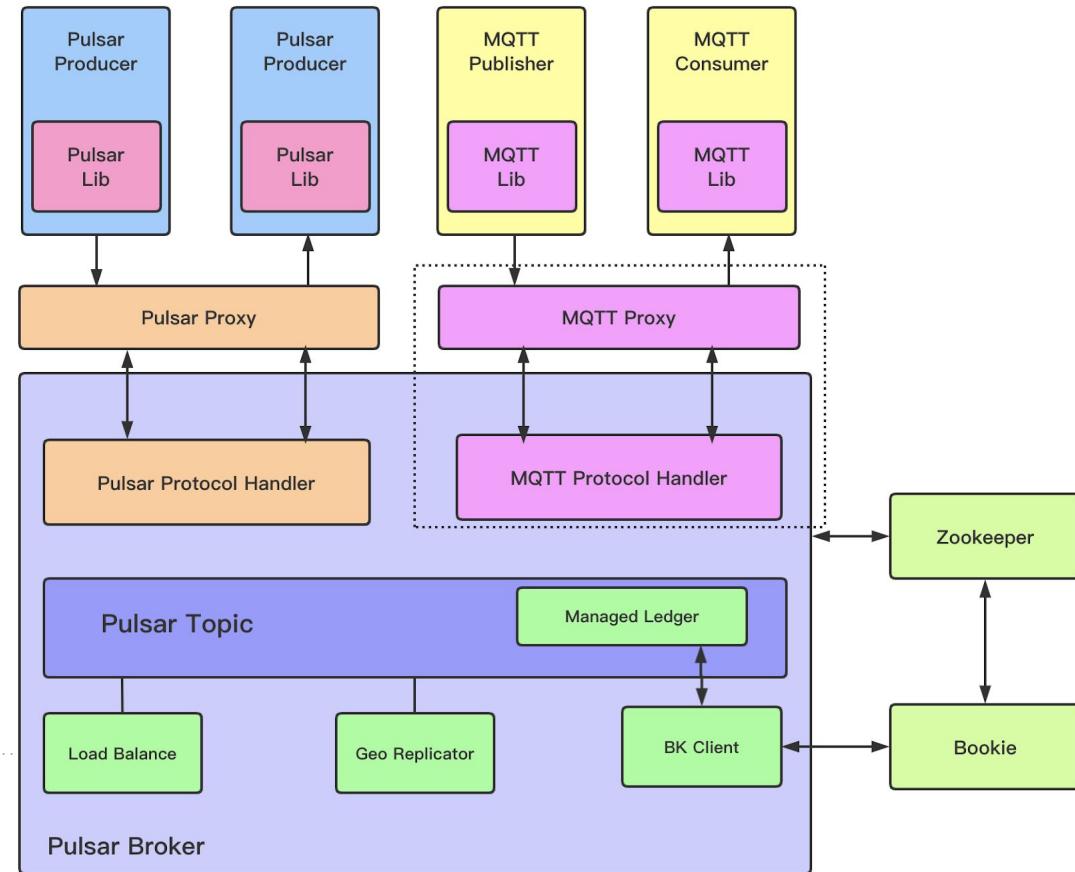
# Moving Data In and Out of Pulsar

IO/Connectors are a simple way to integrate with external systems and move data in and out of Pulsar.

- Built on top of Pulsar Functions
- Built-in connectors - [hub.streamnative.io](https://hub.streamnative.io)

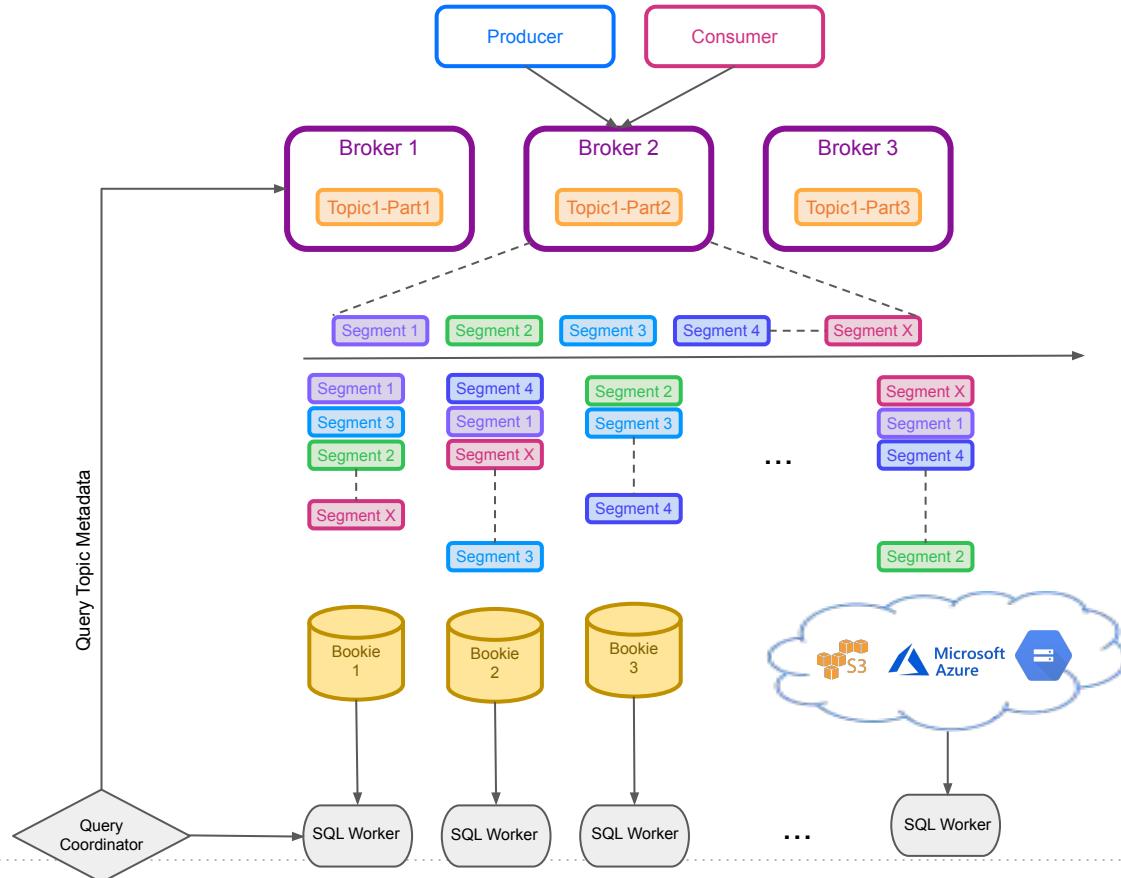


# MQTT on Pulsar (MoP)



# Pulsar SQL

Presto/Trino workers can read segments directly from bookies (or offloaded storage) in parallel.



# Ingesting IoT Data via Java Pulsar

```
UUID uuidKey = UUID.randomUUID();
String pulsarKey = uuidKey.toString();
String OS = System.getProperty("os.name").toLowerCase();
String message = "" + jct.message;
IoTMessage iotMessage = parseMessage("") + jct.message);
String topic = DEFAULT_TOPIC;
if ( jct.topic != null && jct.topic.trim().length()>0) {
    topic = jct.topic.trim();
}
ProducerBuilder<IoTMessage> producerBuilder = client.newProducer(JSONSchema.of(IoTMessage.class))
    .topic(topic)
    .producerName("jetson")
    .sendTimeout(5, TimeUnit.SECONDS);

Producer<IoTMessage> producer = producerBuilder.create();

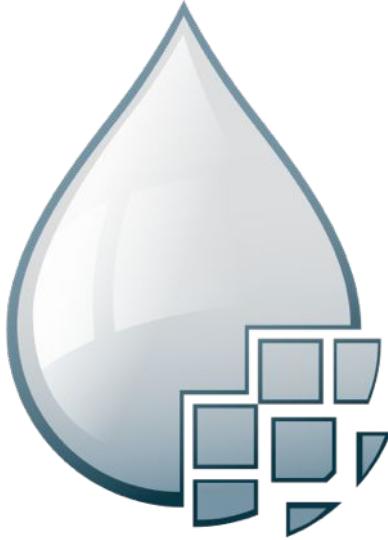
MessageId msgID = producer.newMessage()
    .key(iotMessage.getUuid())
    .value(iotMessage)
    .send();
```

<https://github.com/tspannhw/StreamingAnalyticsUsingFlinkSQL/>

# Ingesting IoT Data via Java Pulsar

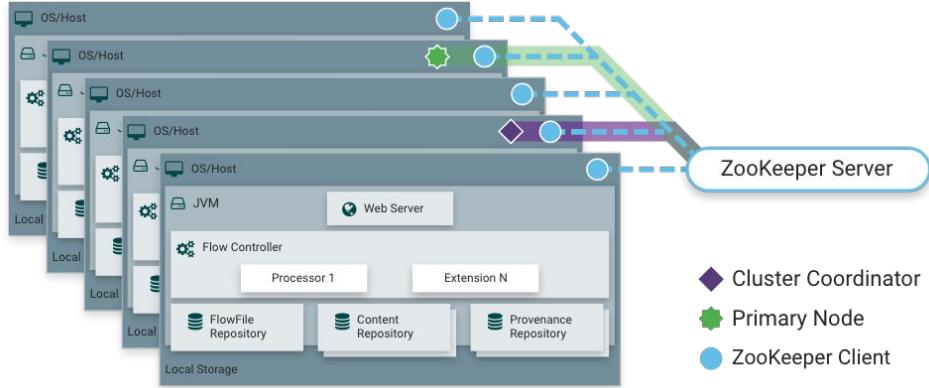
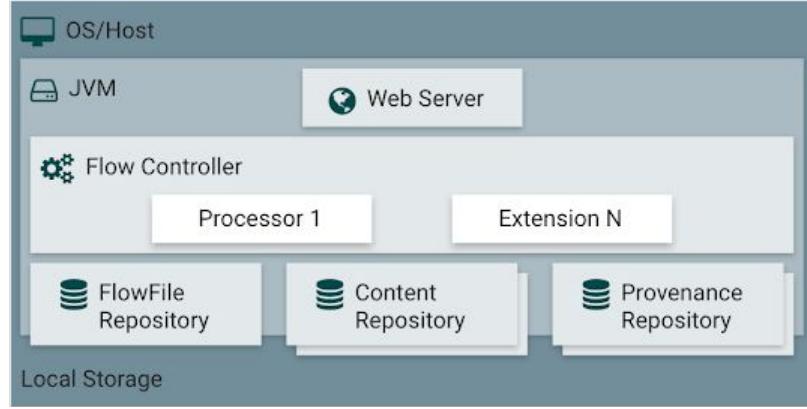
```
private static IoTMessage parseMessage(String message) {  
    IoTMessage iotMessage = null;  
  
    try {  
        if ( message != null && message.trim().length() > 0 ) {  
            ObjectMapper mapper = new ObjectMapper();  
            iotMessage = mapper.readValue(message, IoTMessage.class);  
            mapper = null;  
        }  
    }  
    catch(Throwable t) {  
        t.printStackTrace();  
    }  
  
    if (iotMessage == null) {  
        iotMessage = new IoTMessage();  
    }  
    return iotMessage;  
}
```

# Why Apache NiFi?



- Guaranteed delivery
- Data buffering
  - Backpressure
  - Pressure release
- Prioritized queuing
- Flow specific QoS
  - Latency vs. throughput
  - Loss tolerance
- Data provenance
- Supports push and pull models
- Hundreds of processors
- Visual command and control
- Over a sixty sources
- Flow templates
- Pluggable/multi-role security
- Designed for extension
- Clustering
- Version Control

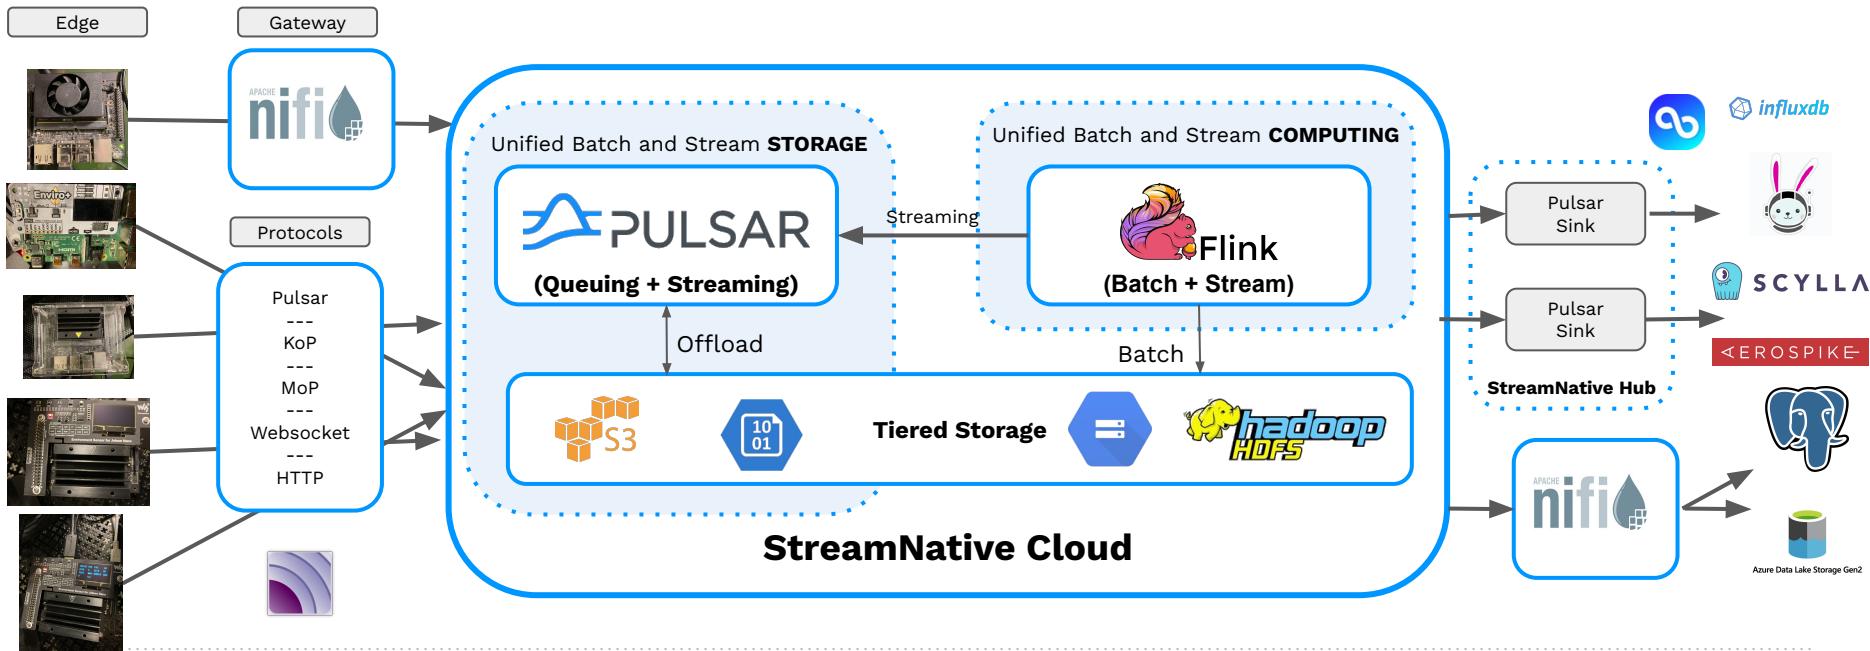
# Architecture



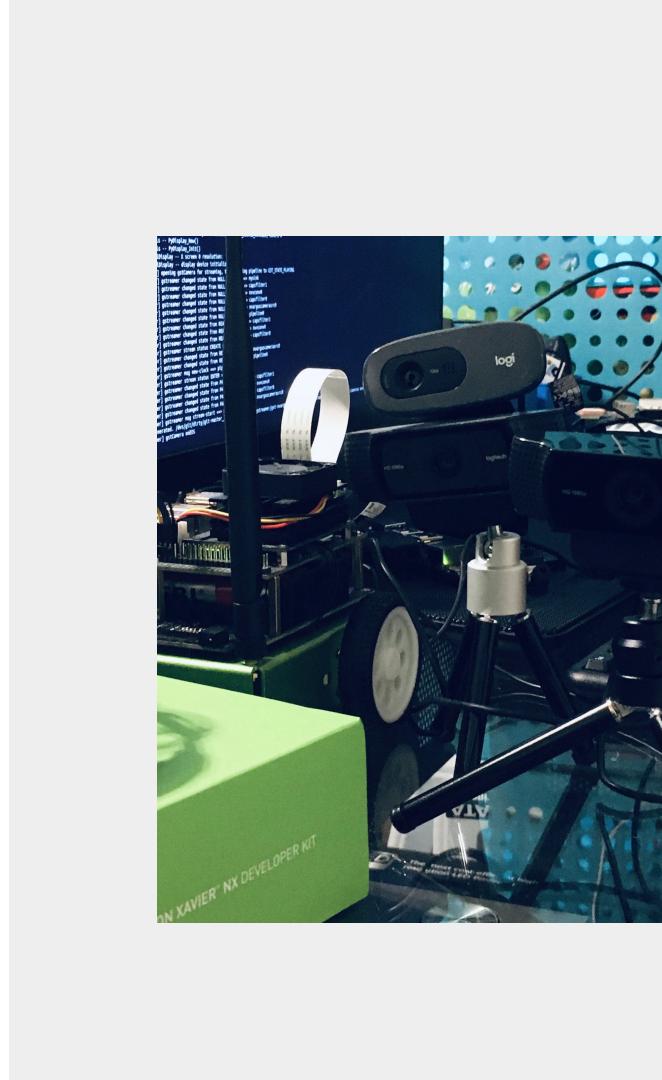
<https://nifi.apache.org/docs/nifi-docs/html/overview.html>

# End-to-End Streaming FLiP(N) IoT Apps

Apache Pulsar - Apache NiFi - MiNiFi <-> Events/Messages <-> Data Stores



# Demo



# Wrap-Up

# Interested In Learning More?



## Resources

[Flink SQL Cookbook](#)

[The Github Source for Flink SQL Demo](#)

[The GitHub Source for Demo](#)



## Free eBooks

[Manning's Apache Pulsar in Action](#)

[O'Reilly Book](#)



## Upcoming Events

[11/8] [PASS Data Community](#)

[11/18] [Developer Week Austin](#)

[11/19] [Porto Tech Hub Con](#)

[12/3] [Data Science Camp](#)

# Let's Keep in Touch!



**Timothy Spann**

Developer Advocate



@PaasDev



<https://www.linkedin.com/in/timothyspann>



<https://github.com/tspannhw>