

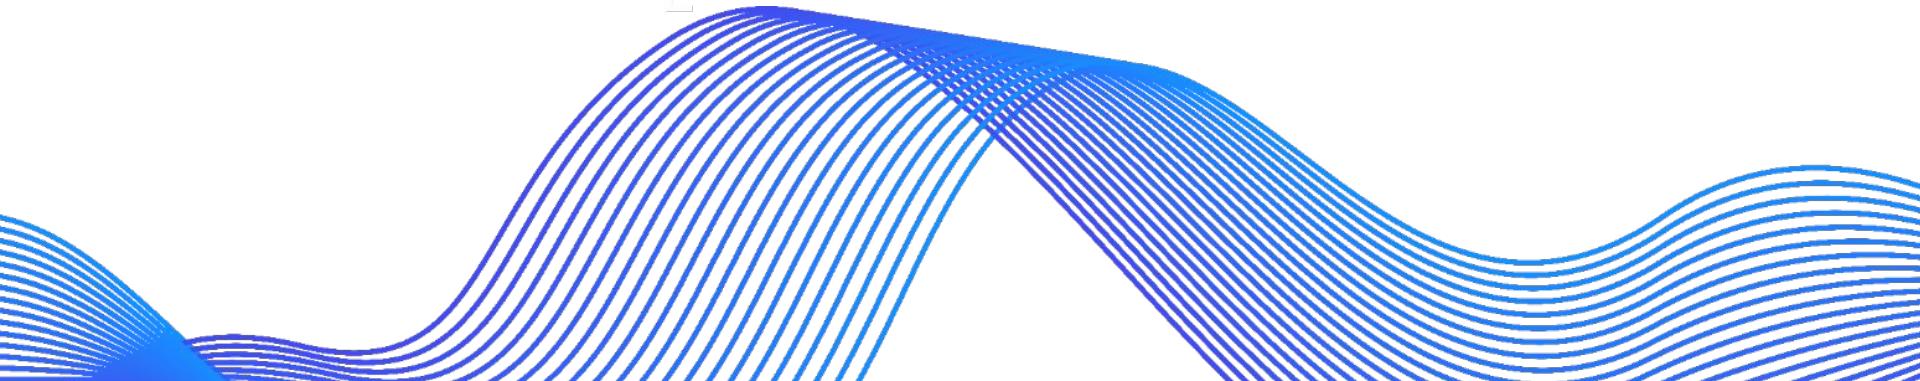


CODE ON
THE BEACH

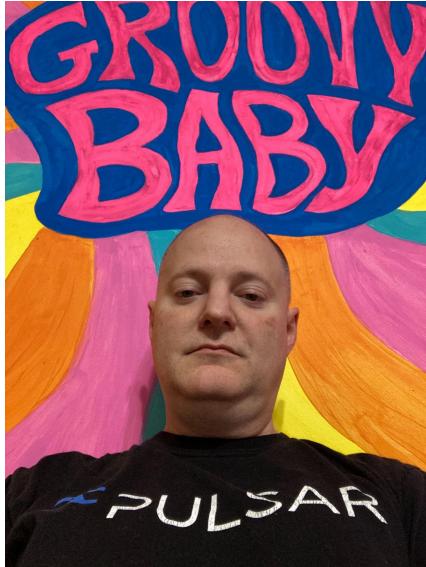
BROUGHT TO
YOU BY
 Availability



Deep Dive into Building Streaming Applications with Apache Pulsar

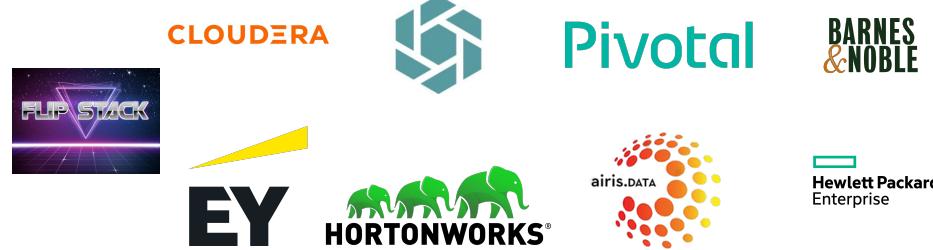






Tim Spann
Developer Advocate

- FLiP(N) Stack = Flink, Pulsar and NiFi Stack
- Streaming Systems/ Data Architect
- Experience:
 - 15+ years of experience with batch and streaming technologies including Pulsar, Flink, Spark, NiFi, Spring, Java, Big Data, Cloud, MXNet, Hadoop, Datalakes, IoT and more.



FLiP Stack Weekly



<https://bit.ly/32dAJft>



This week in Apache Flink, Apache Pulsar, Apache NiFi, Apache Spark and open source friends.

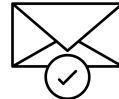


Apache Pulsar is a Cloud-Native
Messaging and Event-Streaming Platform.

Why Apache Pulsar?



Unified
Messaging Platform



Guaranteed
Message Delivery



Resiliency



Infinite
Scalability

Pulsar Benefits



Building
Microservices



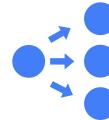
Building Real Time
Applications



Tiered storage



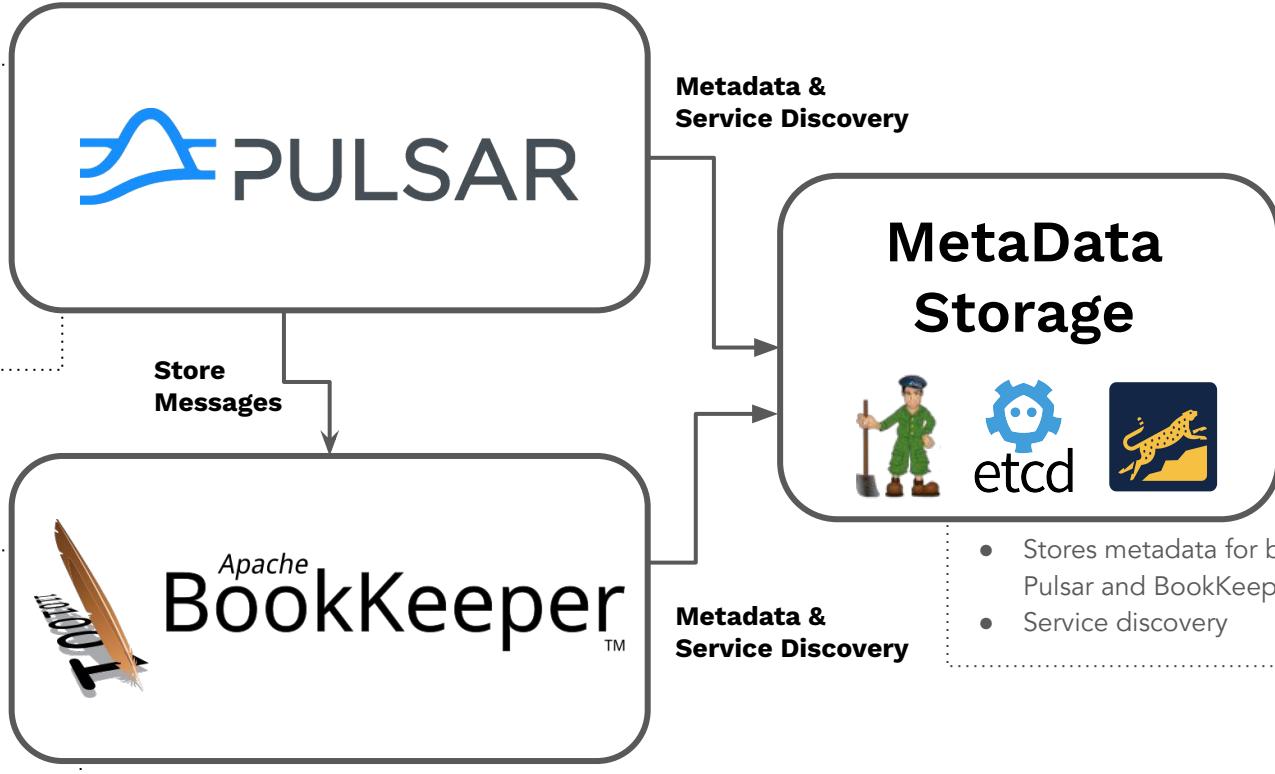
Asynchronous
Communication



Highly Resilient

Key Pulsar Concepts: Architecture

- "Brokers"
- Handles message routing and connections
- Stateless, but with caches
- Automatic load-balancing
- Topics are composed of multiple segments



- "Bookies"
- Stores messages and cursors
- Messages are grouped in segments/ledgers
- A group of bookies form an "ensemble" to store a ledger

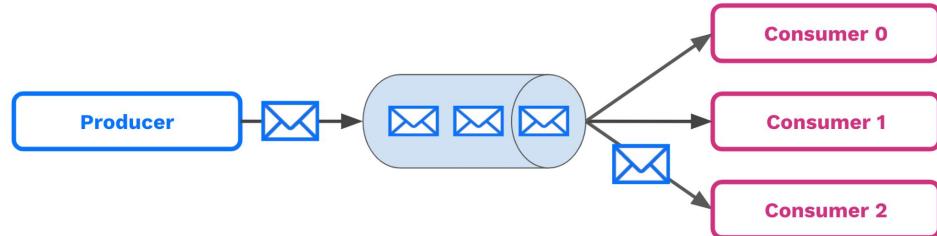
- Stores metadata for both Pulsar and BookKeeper
- Service discovery

Messages - the basic unit of Pulsar

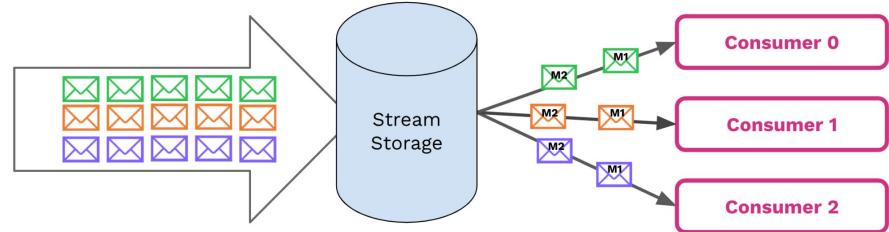
| Component | Description |
|-----------------------------|--|
| Value / data payload | The data carried by the message. All Pulsar messages contain raw bytes, although message data can also conform to data schemas. |
| Key | Messages are optionally tagged with keys, used in partitioning and also is useful for things like topic compaction. |
| Properties | An optional key/value map of user-defined properties. |
| Producer name | The name of the producer who produces the message. If you do not specify a producer name, the default name is used. Message De-Duplication. |
| Sequence ID | Each Pulsar message belongs to an ordered sequence on its topic. The sequence ID of the message is its order in that sequence. Message De-Duplication. |

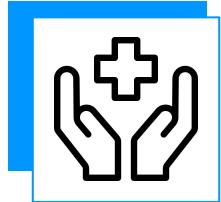
Key Pulsar Concepts: Messaging vs Streaming

Message Queueing - Queueing systems are ideal for work queues that do not require tasks to be performed in a particular order.

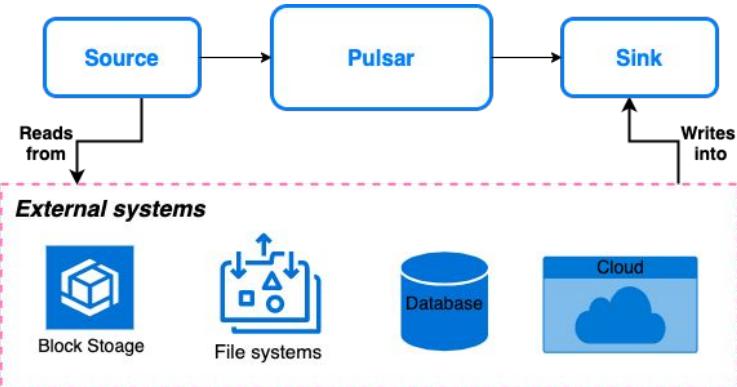


Streaming - Streaming works best in situations where the order of messages is important.





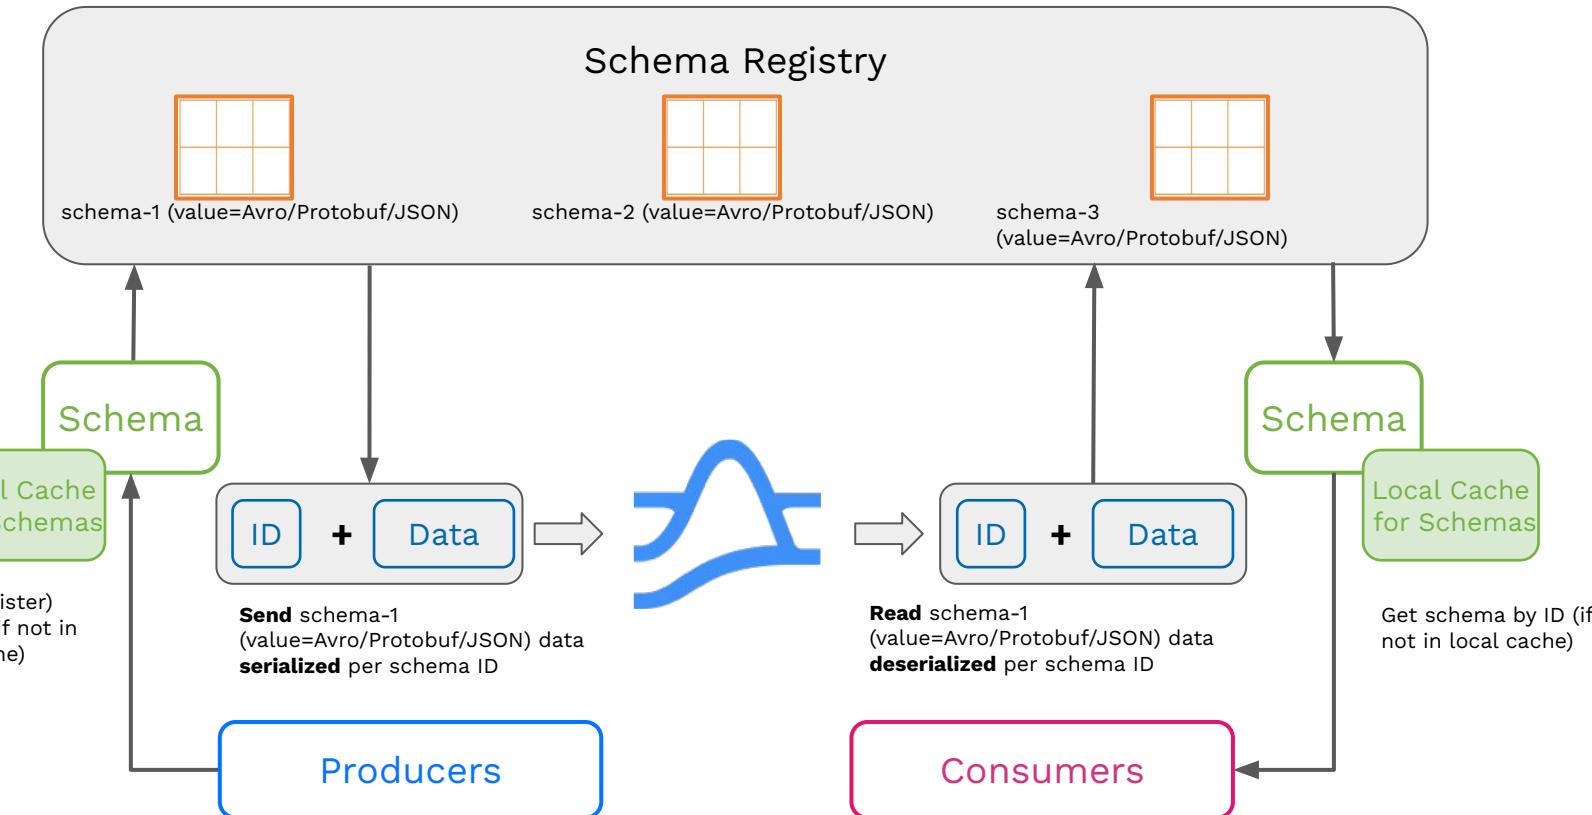
Connectivity



- **Functions** - Lightweight Stream Processing (Java, Python, Go)
- **Connectors** - Sources & Sinks (Cassandra, Kafka, ...)
- **Protocol Handlers** - AoP (AMQP), KoP (Kafka), MoP (MQTT)
- **Processing Engines** - Flink, Spark, Presto/Trino via Pulsar SQL
- **Data Offloaders** - Tiered Storage - (S3)

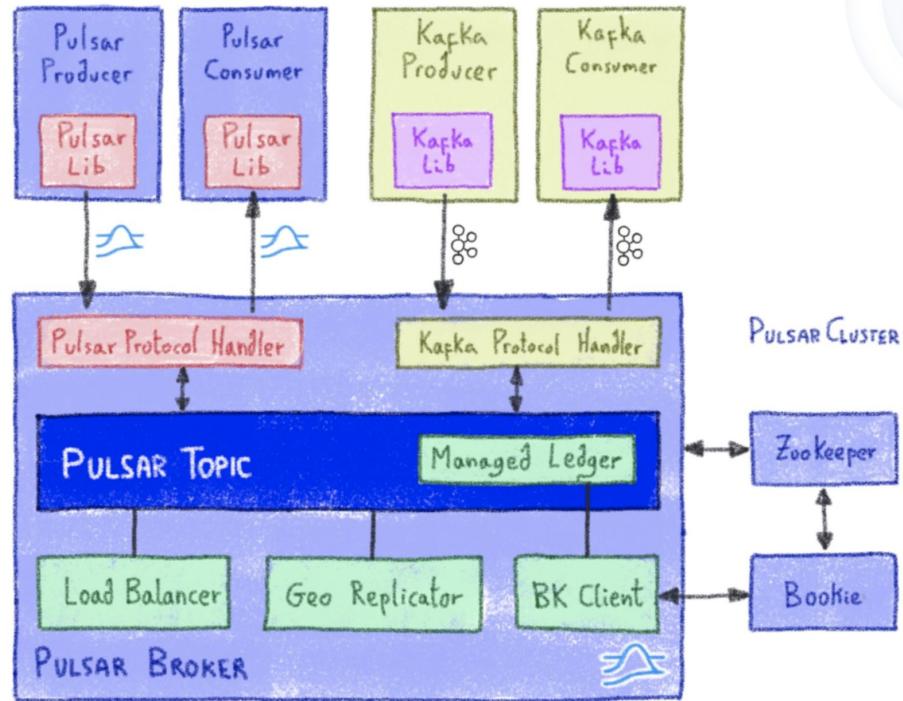
hub.streamnative.io

Schema Registry

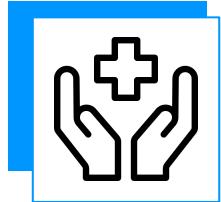




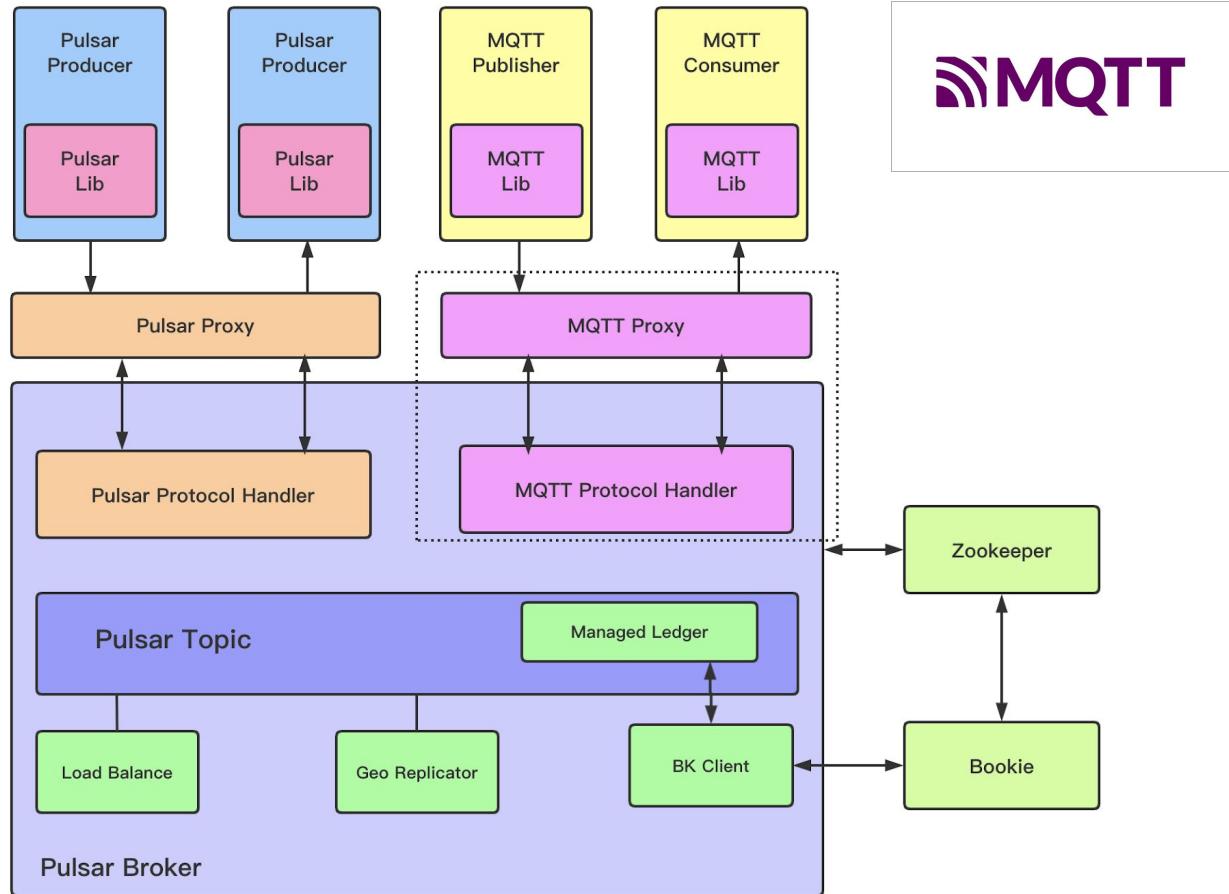
Kafka On Pulsar (KoP)

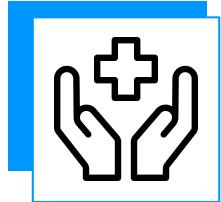


StreamNative

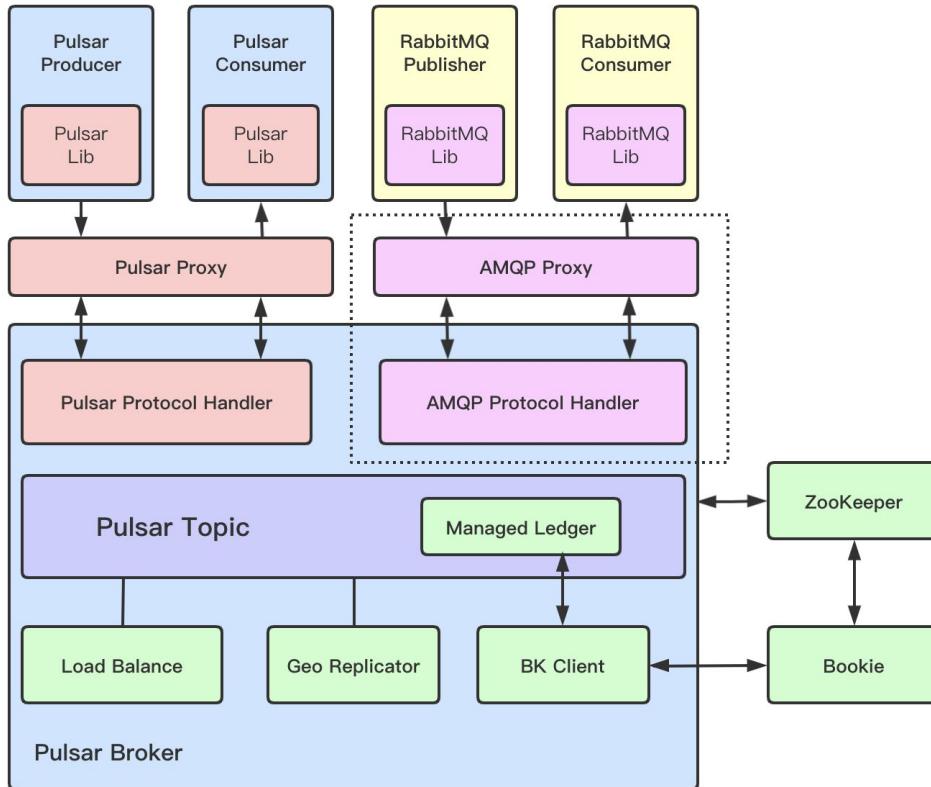


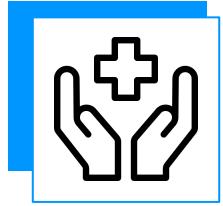
MQTT On Pulsar (MoP)





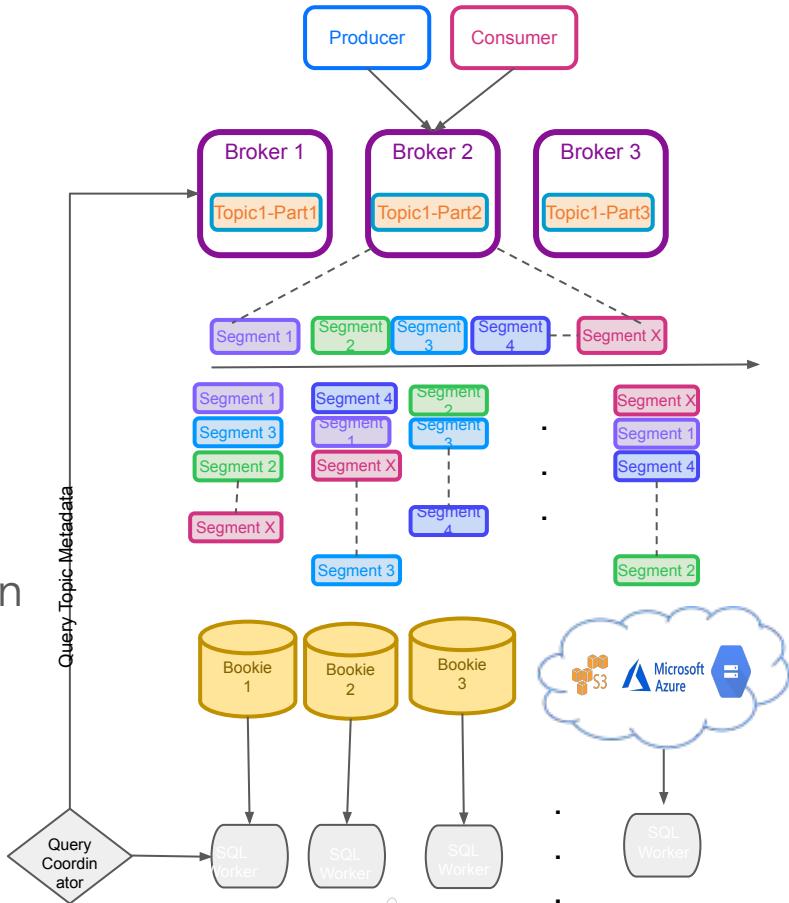
AMQP On Pulsar (AoP)



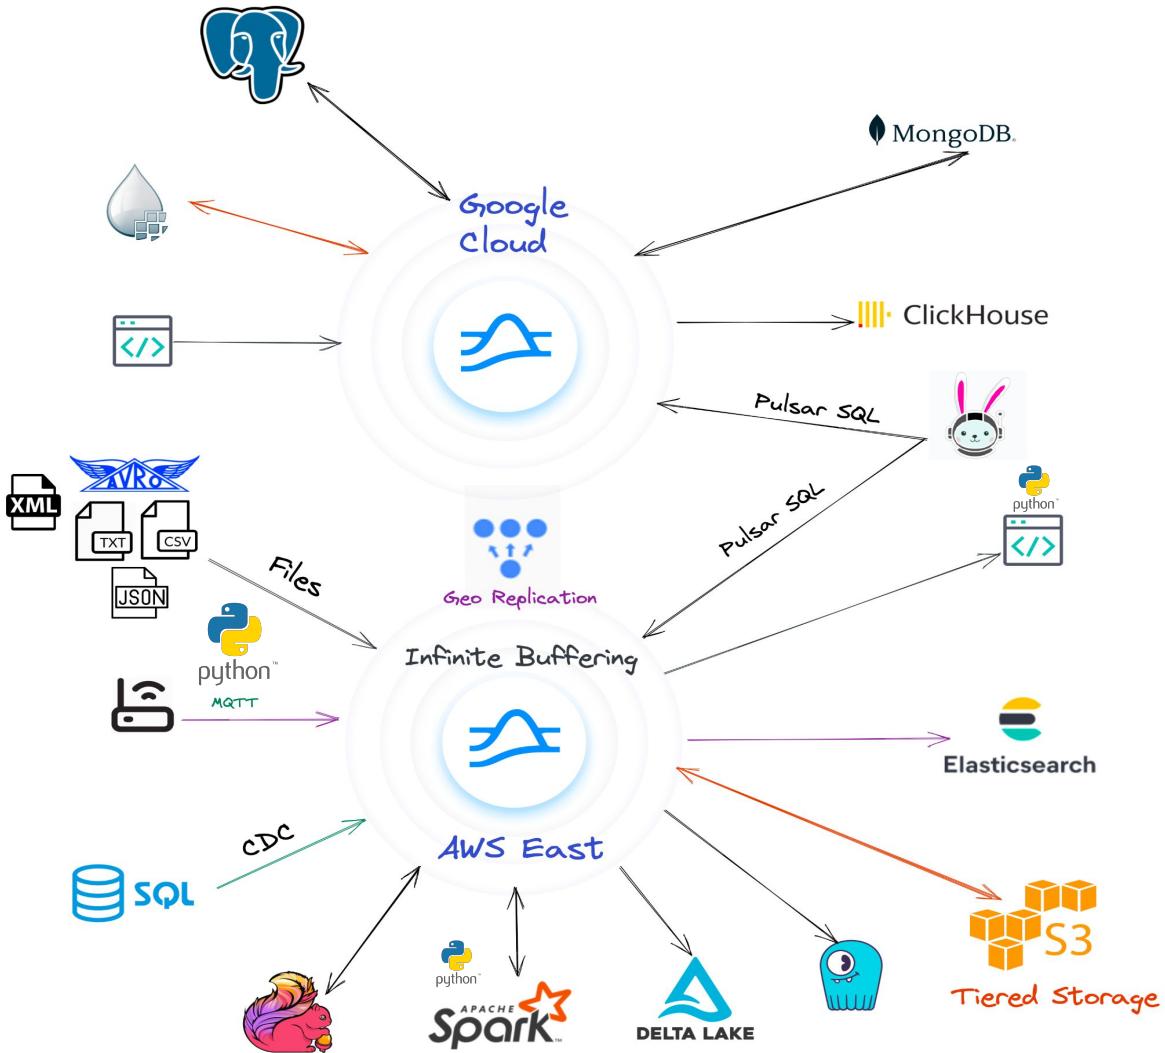


Pulsar SQL

Presto/Trino workers can read segments directly from bookies (or offloaded storage) in parallel.



- Buffer
- Batch
- Route
- Filter
- Aggregate
- Enrich
- Replicate
- Dedupe
- Decouple
- Distribute



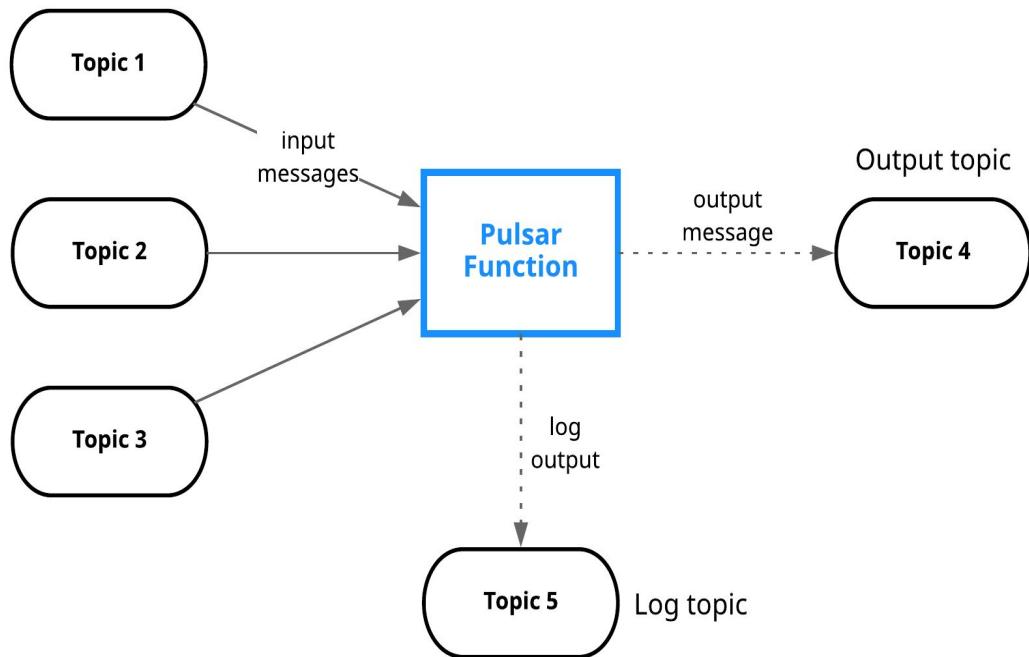
Pulsar Functions

A serverless event streaming framework

- Lightweight computation similar to AWS Lambda.
- Specifically designed to use Apache Pulsar as a message bus.
- Function runtime can be located within Pulsar Broker.

Pulsar Functions

Input topics



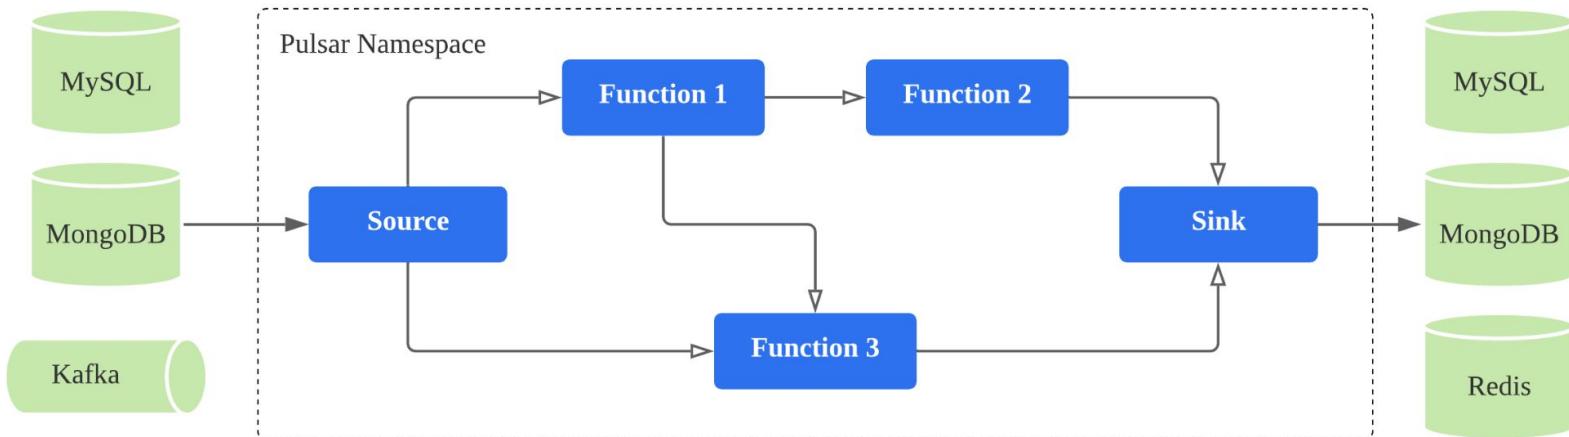
- Consume messages from one or more Pulsar topics.
- Apply user-supplied processing logic to each message.
- Publish the results of the computation to another topic.
- Support multiple programming languages (Java, Python, Go)
- Can leverage 3rd-party libraries to support the ***execution of ML models on the edge.***

Function Mesh

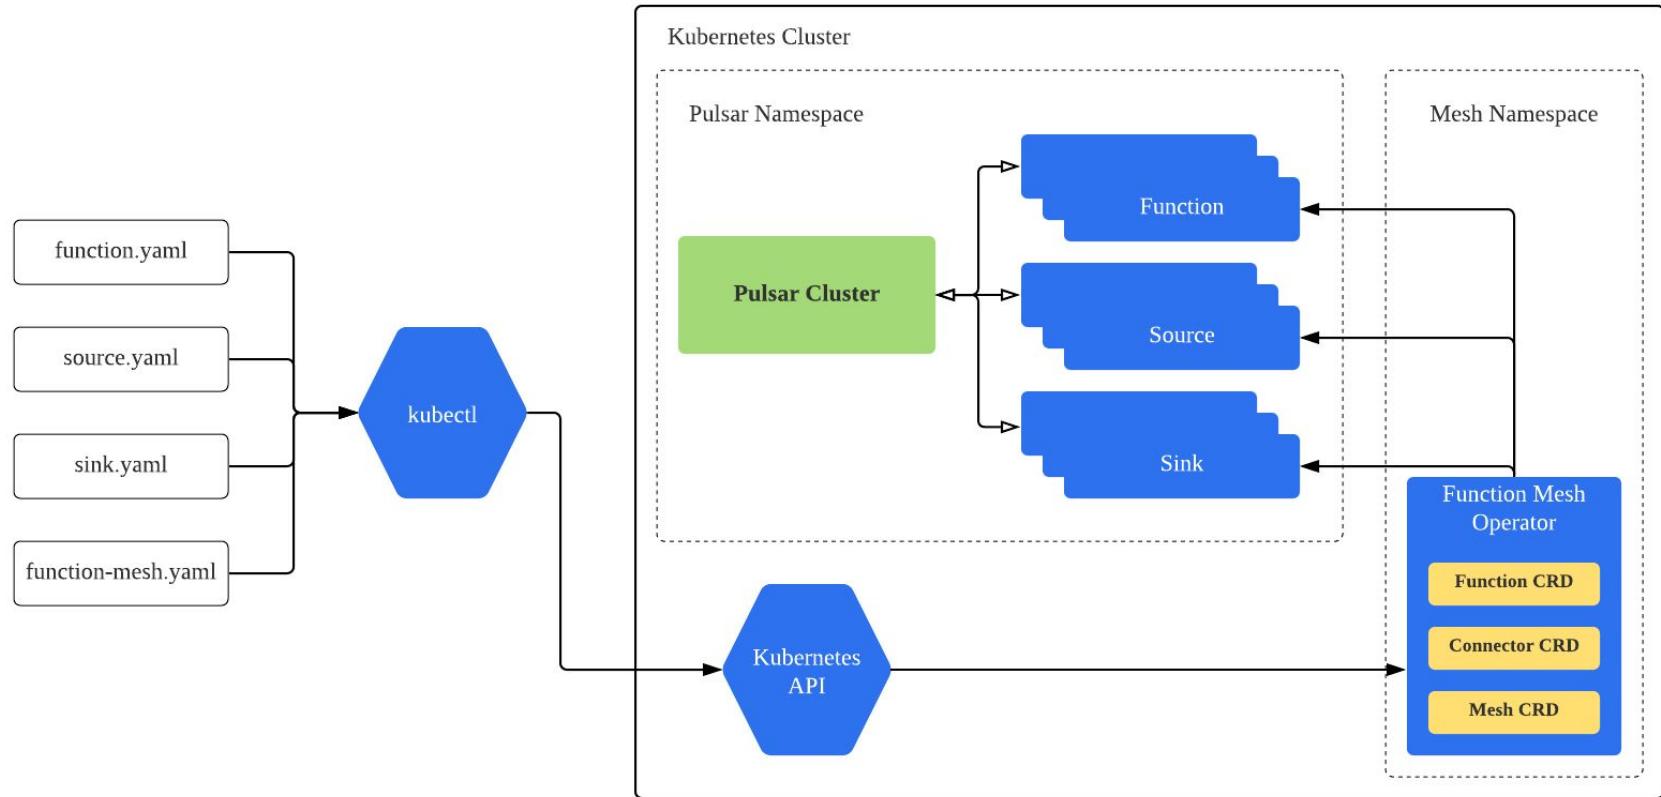


Pulsar Functions, along with Pulsar IO/Connectors, provide a powerful API for ingesting, transforming, and outputting data.

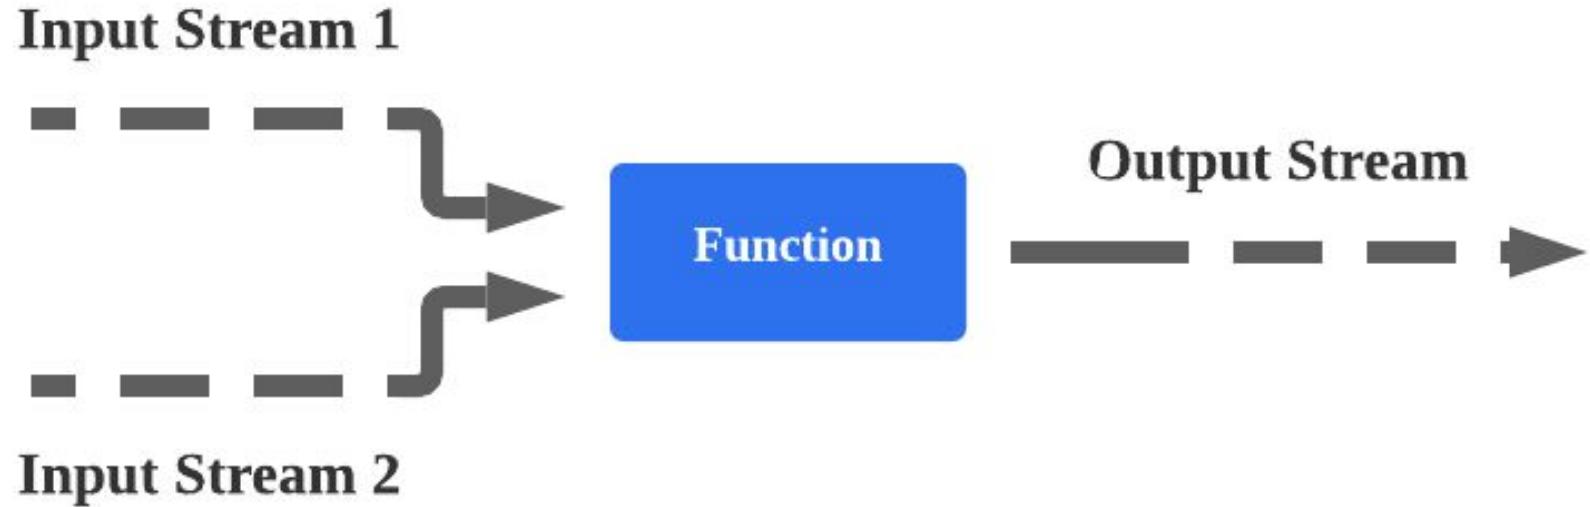
Function Mesh, another StreamNative project, makes it easier for developers to create entire applications built from sources, functions, and sinks all through a declarative API.



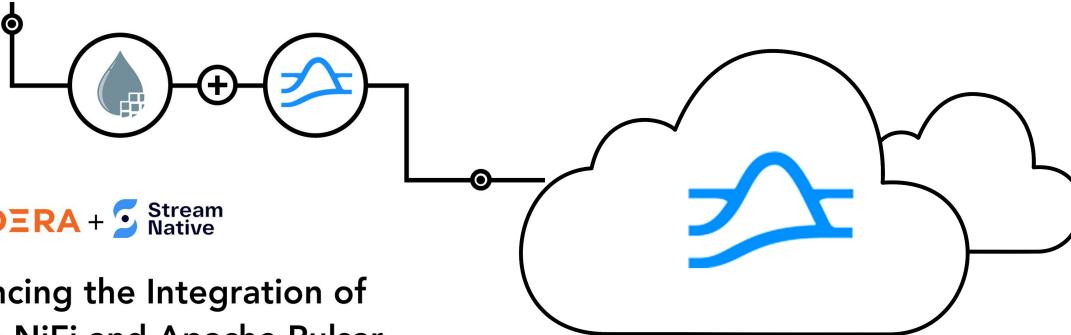
K8 Deploy



Function Execution

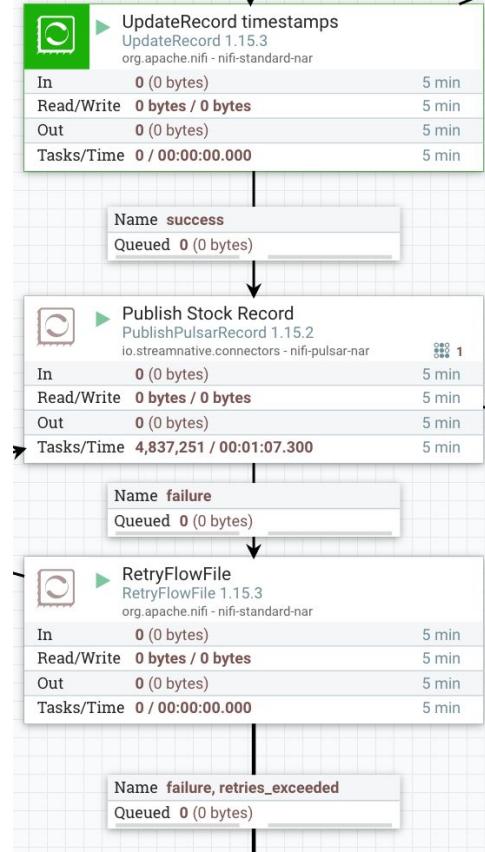


Apache NiFi Pulsar Connector



Announcing the Integration of
Apache NiFi and Apache Pulsar

<https://streamnative.io/apache-nifi-connector/>



Python 3 Coding

Code Along With Tim
=><<DEMO>>



Building a Python3 Producer

```
import pulsar

client = pulsar.Client('pulsar://localhost:6650')
producer
client.create_producer('persistent://conf/ete/first')
producer.send(('Simple Text Message').encode('utf-8'))
client.close()
```

Building a Python 3 Cloud Producer Oauth

```
python3 prod.py -su pulsar+ssl://name1.name2.snio.cloud:6651 -t
persistent://public/default/pyth --auth-params
'{"issuer_url": "https://auth.streamnative.cloud", "private_key": "my.json",
"audience": "urn:sn:pulsar:name:myclustr"}'

from pulsar import Client, AuthenticationOauth2
parse = argparse.ArgumentParser(prog=prod.py')
parse.add_argument('-su', '--service-url', dest='service_url', type=str,
required=True)
args = parse.parse_args()
client = pulsar.Client(args.service_url,
                        authentication=AuthenticationOauth2(args.auth_params))
```

<https://github.com/streamnative/examples/blob/master/cloud/python/OAuth2Producer.py>

Example Avro Schema Usage

```
import pulsar
from pulsar.schema import *
from pulsar.schema import AvroSchema
class thermal(Record):
    uuid = String()
client = pulsar.Client('pulsar://pulsar1:6650')
thermalschema = AvroSchema(thermal)
producer =
    client.create_producer(topic='persistent://public/default/pi-thermal-avro',
                           schema=thermalschema, properties={"producer-name": "thrm"})
thermalRec = thermal()
thermalRec.uuid = "unique-name"
producer.send(thermalRec, partition_key=uniqueid)
```

<https://github.com/tspannhw/FLiP-Pi-Thermal>

Example Json Schema Usage

```
import pulsar
from pulsar.schema import *
from pulsar.schema import JsonSchema
class weather(Record):
    uuid = String()
client = pulsar.Client('pulsar://pulsar1:6650')
wsc = JsonSchema(thermal)
producer =
client.create_producer(topic='persistent://public/default/wthrv', schema=wsc, properties={"producer-name": "wthrv"})
weatherRec = weather()
weatherRec.uuid = "unique-name"
producer.send(weatherRec, partition_key=uniqueid)
```

<https://github.com/tspannhw/FLiP-PulsarDevPython101>

<https://github.com/tspannhw/FLiP-Pi-Weather>

Building a Python3 Consumer

```
import pulsar
client = pulsar.Client('pulsar://localhost:6650')
consumer =
client.subscribe('persistent://conf/ete/first', subscription_name='mine')

while True:
    msg = consumer.receive()
    print("Received message: '%s'" % msg.data())
    consumer.acknowledge(msg)
client.close()
```

MQTT from Python

```
pip3 install paho-mqtt
```

```
import paho.mqtt.client as mqtt
client = mqtt.Client("rpi4-iot")
row = { }
row['gasKO'] = str(readings)
json_string = json.dumps(row)
json_string = json_string.strip()
client.connect("pulsar-server.com", 1883, 180)
client.publish("persistent://public/default/mqtt-2",
payload=json_string,qos=0,retain=True)
```

<https://www.slideshare.net/bunkertor/data-minutes-2-apache-pulsar-with-mqtt-for-edge-computing-lightning-2022>

Web Sockets from Python

```
pip3 install websocket-client
```

```
import websocket, base64, json
topic = 'ws://server:8080/ws/v2/producer/persistent/public/default/topic1'
ws = websocket.create_connection(topic)
message = "Hello Philly ETE Conference"
message_bytes = message.encode('ascii')
base64_bytes = base64.b64encode(message_bytes)
base64_message = base64_bytes.decode('ascii')
ws.send(json.dumps({'payload' : base64_message, 'properties': {'device' : 'macbook'}, 'context' : 5}))
response = json.loads(ws.recv())
```

<https://github.com/tspannhw/FLiP-IoT/blob/main/wsreader.py>

<https://github.com/tspannhw/FLiP-IoT/blob/main/wspulsar.py>

<https://pulsar.apache.org/docs/en/client-libraries-websocket/>

Kafka from Python

```
pip3 install kafka-python

from kafka import KafkaProducer
from kafka.errors import KafkaError

row = { }
row['gasKO'] = str(readings)
json_string = json.dumps(row)
json_string = json_string.strip()

producer = KafkaProducer(bootstrap_servers='pulsar1:9092', retries=3)
producer.send('topic-kafka-1', json.dumps(row).encode('utf-8'))
producer.flush()
```

<https://docs.streamnative.io/platform/v1.0.0/concepts/kop-concepts>

Deploy Python Functions

```
bin/pulsar-admin functions create --auto-ack true --py py/src/sentiment.py  
--classname "sentiment.Chat" --inputs "persistent://public/default/chat"  
--log-topic "persistent://public/default/logs" --name Chat --output  
"persistent://public/default/chatresult"
```

Pulsar IO Function in Python 3

```
from pulsar import Function
import json

class Chat(Function):
    def __init__(self):
        pass

    def process(self, input, context):
        logger = context.get_logger()

        msg_id = context.get_message_id()
        fields = json.loads(input)
```

Building a Golang Pulsar App

```
go get -u "github.com/apache/pulsar-client-go/pulsar"

import (
    "log"
    "time"
    "github.com/apache/pulsar-client-go/pulsar"
)
func main() {
    client, err := pulsar.NewClient(pulsar.ClientOptions{
        URL: "pulsar://localhost:6650", OperationTimeout: 30 * time.Second,
        ConnectionTimeout: 30 * time.Second,
    })
    if err != nil {
        log.Fatalf("Could not instantiate Pulsar client: %v", err)
    }
    defer client.Close()
}
```



Typed Java Client

```
Producer<User> producer = client.newProducer(Schema.AVRO(User.class)).create();
producer.newMessage()
    .value(User.builder()
        .userName("pulsar-user")
        .userId(1L)
        .build())
    .send();
```

```
Consumer<User> consumer =
client.newConsumer(Schema.AVRO(User.class)).create();
User user = consumer.receive();
```

Using Connection URLs

Java:

```
PulsarClient client = PulsarClient.builder()  
    .serviceUrl("pulsar://broker1:6650")  
    .build();
```

C#:

```
var client = new PulsarClientBuilder()  
    .ServiceUrl("pulsar://broker1:6650")  
    .Build();
```

Auth

C#

```
// Create the auth class
var tokenAuth = AuthenticationFactory.token("<jwttoken>");

// Pass the authentication implementation
var client = new PulsarClientBuilder()
    .ServiceUrl("pulsar://broker1:6650")
    .Authentication(tokenAuth)
    .Build();
```

Source Code

<https://github.com/tspannhw/airquality>

<https://github.com/tspannhw/FLiPN-AirQuality-REST>

<https://github.com/tspannhw/pulsar-airquality-function>

<https://github.com/tspannhw/FLiP-Pi-BreakoutGarden>

<https://github.com/tspannhw/FLiPN-DEVNEXUS-2022>

Pulsar

Function Java

Your Code Here



```
import java.util.function.Function;  
  
public class MyFunction implements Function<String, String> {  
    public String apply(String input) {  
        return doBusinessLogic(input);  
    }  
}
```

The incoming messages are passed
into the function one-by-one

The returned value is automatically
published to the output topic

Pulsar

Function SDK

Your Code Here



```
import org.apache.pulsar.client.impl.schema.JSONSchema;
import org.apache.pulsar.functions.api.*;

public class AirQualityFunction implements Function<byte[], Void> {
    @Override
    public Void process(byte[] input, Context context) {
        context.getLogger().debug("File:" + new String(input));
        context.newOutputMessage("topicname",
            JSONSchema.of(Observation.class))
            .key(UUID.randomUUID().toString())
            .property("prop1", "value1")
            .value(observation)
            .send();
    }
}
```

Setting Subscription Type Java

```
Consumer<byte[]> consumer = pulsarClient.newConsumer()  
    .topic(topic)  
    .subscriptionName("subscriptionName")  
    .subscriptionType(SubscriptionType.Shared)  
    .subscribe();
```

Subscribing to a Topic and setting Subscription Name Java

```
Consumer<byte[]> consumer = pulsarClient.newConsumer()  
    .topic(topic)  
    .subscriptionName("subscriptionName")  
    .subscribe();
```

Producing Object Events From Java

```
ProducerBuilder<Observation> producerBuilder =  
pulsarClient.newProducer(JSONSchema.of(Observation.class))  
    .topic(topicName)  
    .producerName(producerName).sendTimeout(60,  
                                              TimeUnit.SECONDS);  
Producer<Observation> producer = producerBuilder.create();
```

```
msgID = producer.newMessage()  
    .key(someUniqueKey)  
    .value(observation)  
    .send();
```

Building Pulsar SQL View

```
bin/pulsar sql
```

```
show catalogs;  
  
show schemas in pulsar;  
  
show tables in pulsar."conf/ete";  
  
select * from pulsar."conf/ete"."first";  
  
exit;
```

Spark + Pulsar



```
val dfPulsar = spark.readStream.format("pulsar")
    .option("service.url", "pulsar://pulsar1:6650")
    .option("admin.url", "http://pulsar1:8080")
    .option("topic", "persistent://public/default/airquality") .load()
```

```
val pQuery = dfPulsar.selectExpr("*")
.writeStream.format("console")
.option("truncate", false).start()
```

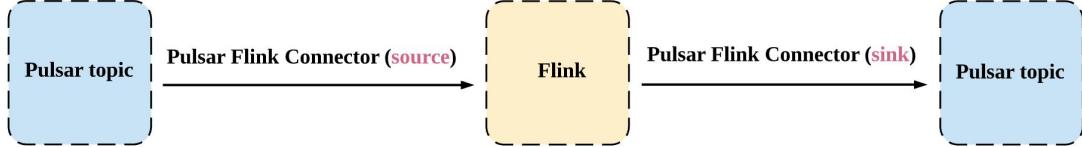
<https://pulsar.apache.org/docs/en/adaptors-spark/>



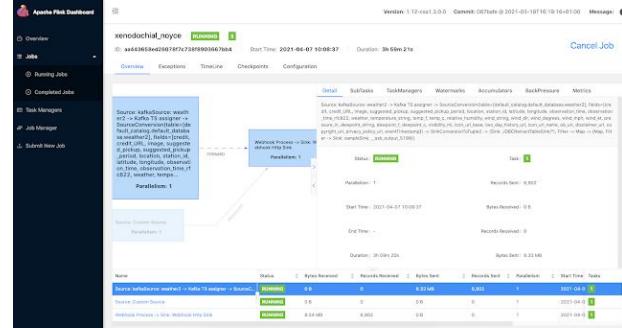
version 3.2.0

Using Scala version 2.12.15
(OpenJDK 64-Bit Server VM, Java 11.0.11)

Apache Flink?



- Unified computing engine
- Batch processing is a special case of stream processing
- Stateful processing
- Massive Scalability
- Flink SQL for queries, inserts against Pulsar Topics
- Streaming Analytics
- Continuous SQL
- Continuous ETL
- Complex Event Processing
- Standard SQL Powered by Apache Calcite



Building Pulsar Flink SQL View

```
CREATE CATALOG pulsar WITH (
    'type' = 'pulsar',
    'service-url' = 'pulsar://pulsar1:6650',
    'admin-url' = 'http://pulsar1:8080',
    'format' = 'json'
);

select * from anytopicisatable;
```



Flink SQL

```
select aqi, parameterName, dateObserved, hourObserved, latitude,  
longitude, localTimeZone, stateCode, reportingArea from  
airquality;
```

```
select max(aqi) as MaxAQI, parameterName, reportingArea from  
airquality group by parameterName, reportingArea;
```

```
select max(aqi) as MaxAQI, min(aqi) as MinAQI, avg(aqi) as  
AvgAQI, count(aqi) as RowCount, parameterName, reportingArea  
from airquality group by parameterName, reportingArea;
```

Welcome! Enter 'HELP;' to list all available commands. 'QUIT;' to exit.

```
Flink SQL> CREATE CATALOG pulsar WITH (  
>   'type' = 'pulsar',  
>   'service-url' = 'pulsar://pulsar1:6650',  
>   'admin-url' = 'http://pulsar1:8080',  
>   'format' = 'json'  
> );  
[INFO] Execute statement succeed.  
Flink SQL> █
```



Flink SQL

Refresh: 1 s SQL Query Result (Table) Page: Last of 1 Updated: 15:52:48.189

| uuid | ipaddress | cputempf | runtime |
|--------------------|---------------|----------|---------|
| snr_20220323195238 | 192.168.1.229 | 99 | 453 |
| snr_20220323195243 | 192.168.1.229 | 100 | 458 |

I

Q Quit R Refresh + Inc Refresh - Dec Refresh G Goto Page L Last Page N Next Page P Prev Page O Open Row

Building Spark SQL View



```
val dfPulsar = spark.readStream.format("pulsar")
    .option("service.url", "pulsar://pulsar1:6650")
    .option("admin.url", "http://pulsar1:8080")
    .option("topic", "persistent://public/default/pi-sensors")
    .load()

dfPulsar.printSchema()

val pQuery = dfPulsar.selectExpr("*")
    .writeStream.format("console")
    .option("truncate", false)
    .start()
```

Monitoring and Metrics Check

```
curl http://localhost:8080/admin/v2/persistent/conf/ete/first/stats |  
python3 -m json.tool
```

```
bin/pulsar-admin topics stats-internal persistent://conf/ete/first
```

```
curl http://pulsar1:8080/metrics/
```

```
bin/pulsar-admin topics stats-internal persistent://conf/ete/first
```

```
bin/pulsar-admin topics peek-messages --count 5 --subscription ete-reader  
persistent://conf/ete/first
```

```
bin/pulsar-admin topics subscriptions persistent://conf/ete/first
```

Cleanup

```
bin/pulsar-admin topics delete persistent://conf/ete/first
```

```
bin/pulsar-admin namespaces delete conf/ete
```

```
bin/pulsar-admin tenants delete conf
```

Metrics: Broker

Broker metrics are exposed under "`/metrics`" at port **8080**.

You can change the port by updating `webServicePort` to a different port in the `broker.conf` configuration file.

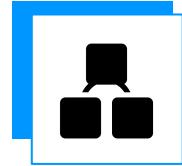
All the metrics exposed by a broker are labeled with `cluster=${pulsar_cluster}`.

The name of Pulsar cluster is the value of `${pulsar_cluster}`, configured in the `broker.conf` file.

For more information: <https://pulsar.apache.org/docs/en/reference-metrics/#broker>

These metrics are available for brokers:

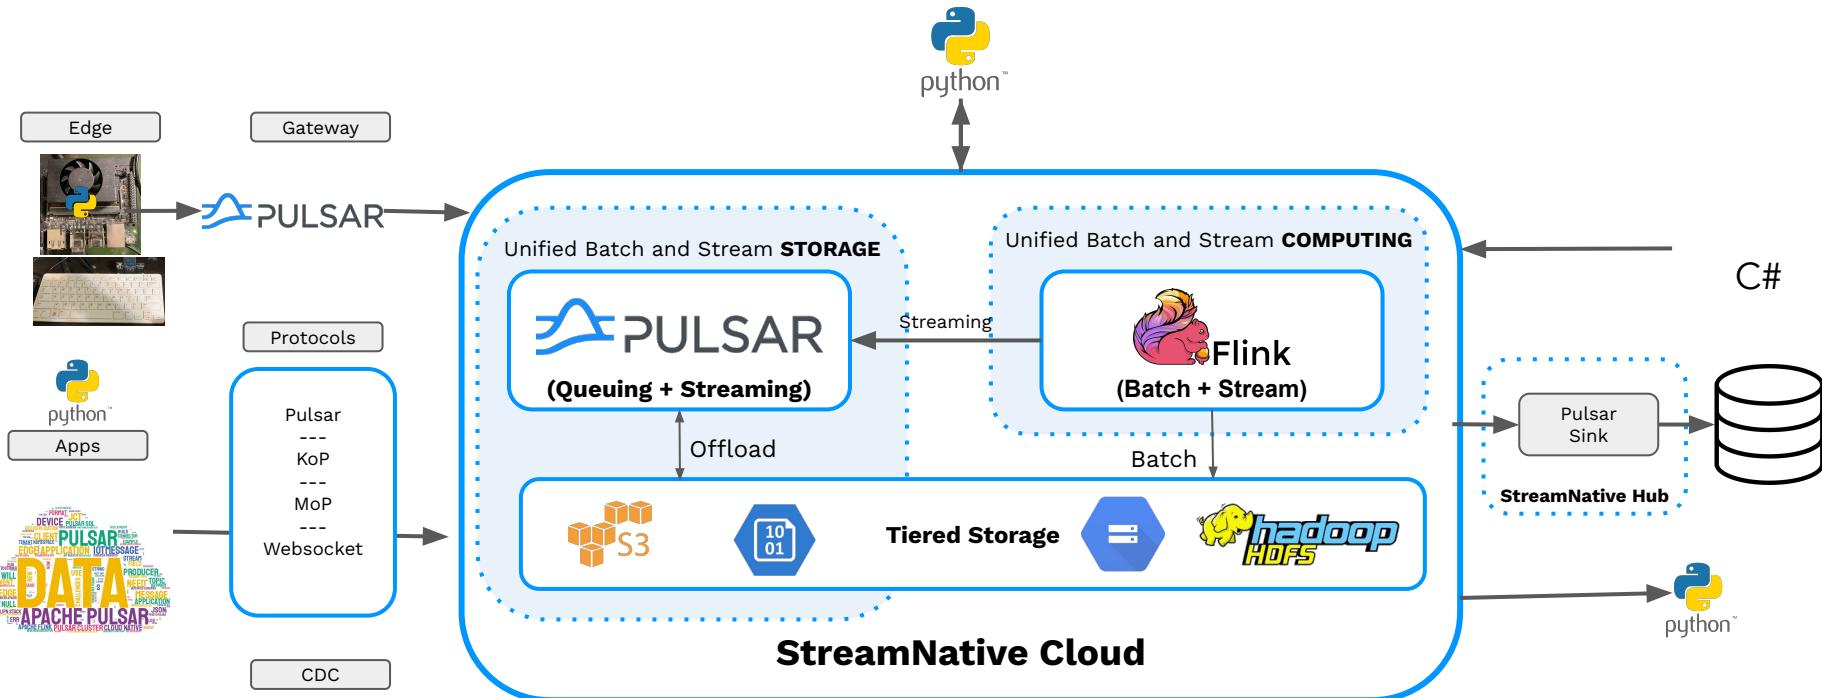
- Namespace metrics
 - Replication metrics
- Topic metrics
 - Replication metrics
- ManagedLedgerCache metrics
- ManagedLedger metrics
- LoadBalancing metrics
 - BundleUnloading metrics
 - BundleSplit metrics
- Subscription metrics
- Consumer metrics
- ManagedLedger bookie client metrics

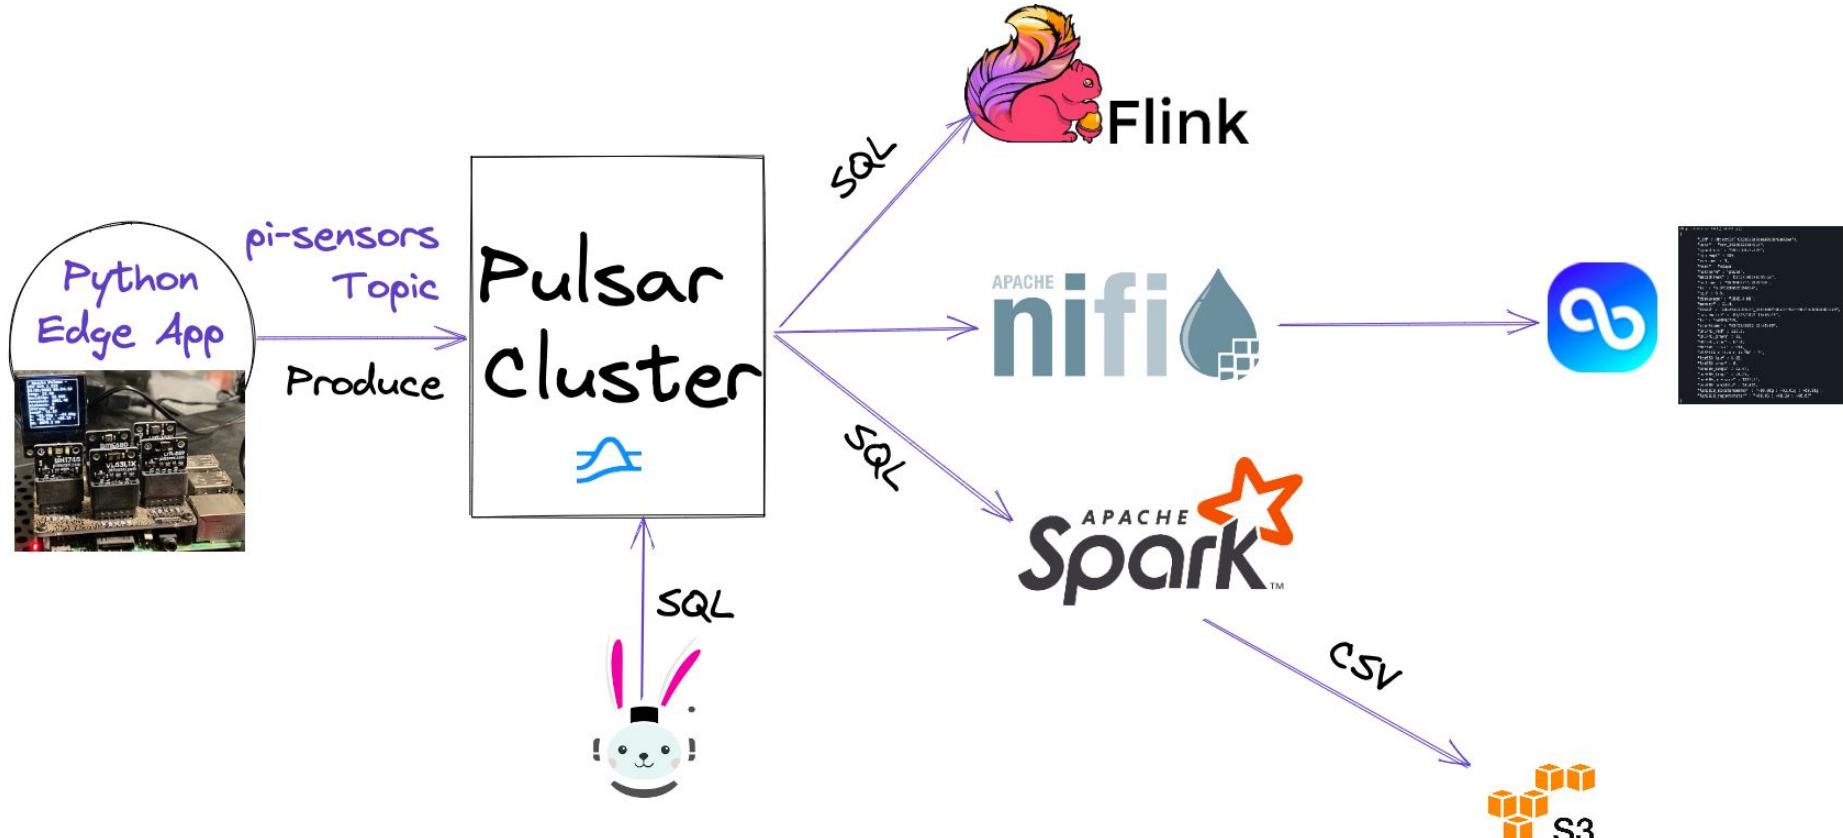


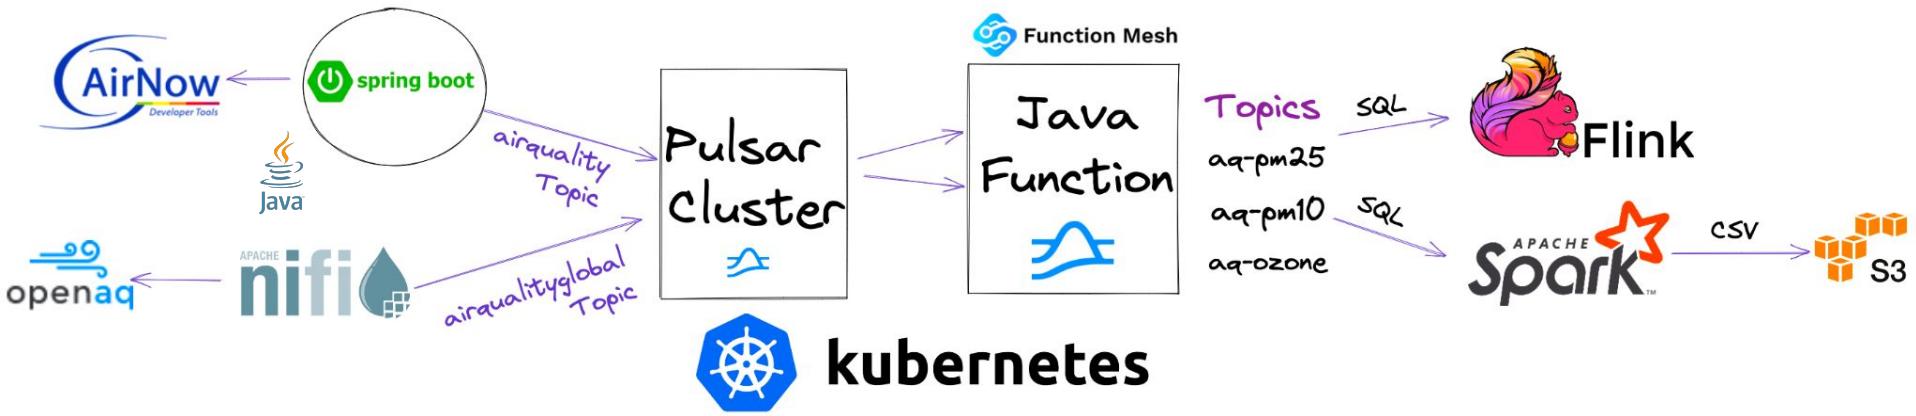
Use Cases

- Unified Messaging Platform
- AdTech
- Fraud Detection
- Connected Car
- IoT Analytics
- Microservices Development

Streaming FLiP-Py Apps







Pulsar Summit

San Francisco

Hotel Nikko

August 18 2022



[Save Your Spot Now](#)

Use code CODEONTHEBEACH20
to get 20% off.

5 Keynotes
12 Breakout Sessions
1 Amazing Happy Hour



databricks



Pulsar Summit San Francisco Sponsorship Prospectus

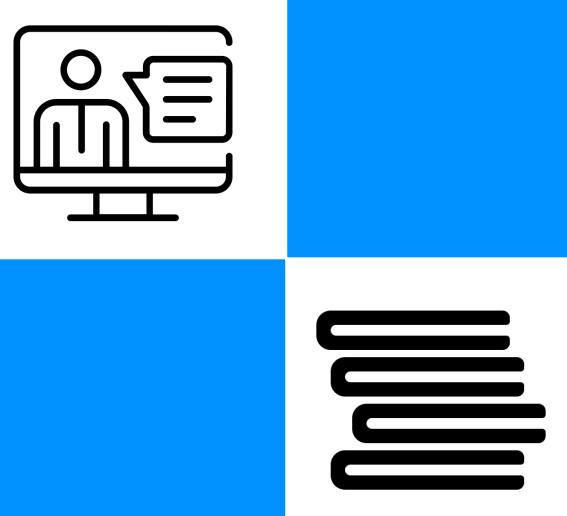


Community Sponsorships Available

Help engage and connect the Apache Pulsar community by becoming an official sponsor for Pulsar Summit San Francisco 2022! Learn more about the requirements and benefits of becoming a community sponsor.

RESOURCES

Here are resources to continue your journey
with Apache Pulsar



Python For Pulsar on Pi

- <https://github.com/tspannhw/FLiP-Pi-BreakoutGarden>
- <https://github.com/tspannhw/FLiP-Pi-Thermal>
- <https://github.com/tspannhw/FLiP-Pi-Weather>
- <https://github.com/tspannhw/FLiP-RP400>
- <https://github.com/tspannhw/FLiP-Py-Pi-GasThermal>
- <https://github.com/tspannhw/FLiP-PY-FakeDataPulsar>
- <https://github.com/tspannhw/FLiP-Py-Pi-EnviroPlus>
- <https://github.com/tspannhw/PythonPulsarExamples>
- <https://github.com/tspannhw/pulsar-pychat-function>
- <https://github.com/tspannhw/FLiP-PulsarDevPython101>
- <https://github.com/tspannhw/airquality>





Founded by the original developers of Apache Pulsar.

Passionate and dedicated team.

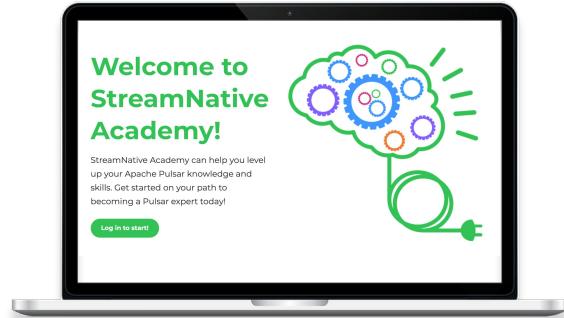
StreamNative helps teams to **capture**, **manage**, and **leverage data** using Pulsar's unified messaging and streaming platform.

streamnative.io

StreamNative Academy

Apache Pulsar Training

- Instructor-led courses
 - Pulsar Fundamentals
 - Pulsar Developers
 - Pulsar Operations
- On-demand learning with labs
- 300+ engineers, admins and architects trained!



**Now Available
On-Demand
Pulsar Training**

Academy.StreamNative.io





Let's Keep in Touch!



Tim Spann

Developer Advocate



[PaaSDev](#)



<https://www.linkedin.com/in/timothyspann>



<https://github.com/tspannhw>