



Let's keep it simple and streaming

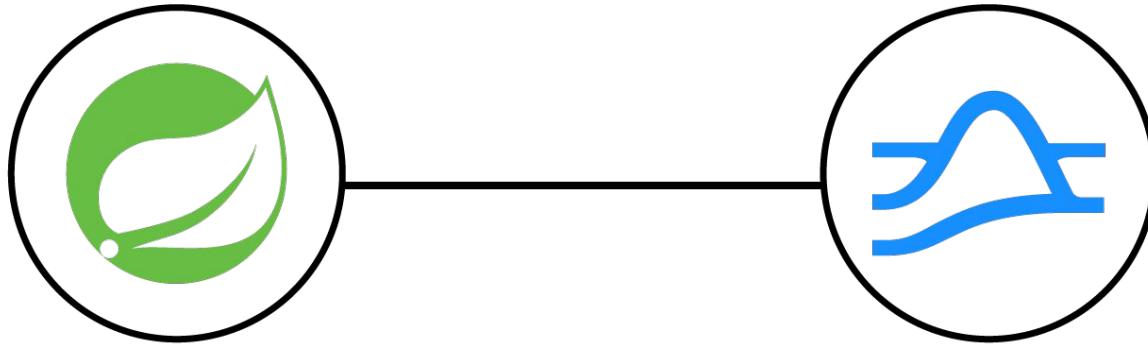
Tim Spann
Developer Advocate

- FLiP(N) Stack = Flink, Pulsar and NiFi Stack
- Streaming Systems/ Data Architect
- Experience:
 - 15+ years of experience with batch and streaming technologies including Pulsar, Flink, Spark, NiFi, Spring, Java, Big Data, Cloud, MXNet, Hadoop, Datalakes, IoT and more.



Tim Spann
Developer Advocate





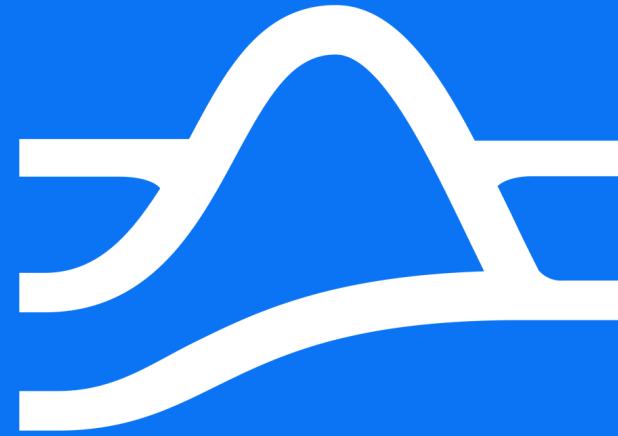
Announcing Spring for Apache Pulsar

APACHE PULSAR

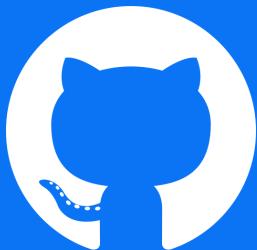
EVERYWHERE

Agenda

- Introduction
- What is Apache Pulsar?
- Spring Apps
- Pulsar
- AMQP
- MQTT
- Kafka
- Demo

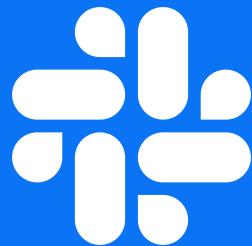


Apache Pulsar has a vibrant community



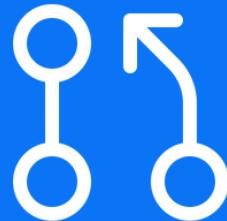
560+

Contributors



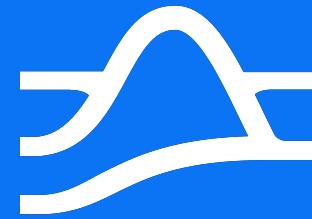
7,000+

Slack Members



10,000+

Commits



1,000+

Organizations
Using Pulsar



PULSAR 101



**Unified
Messaging
Platform**



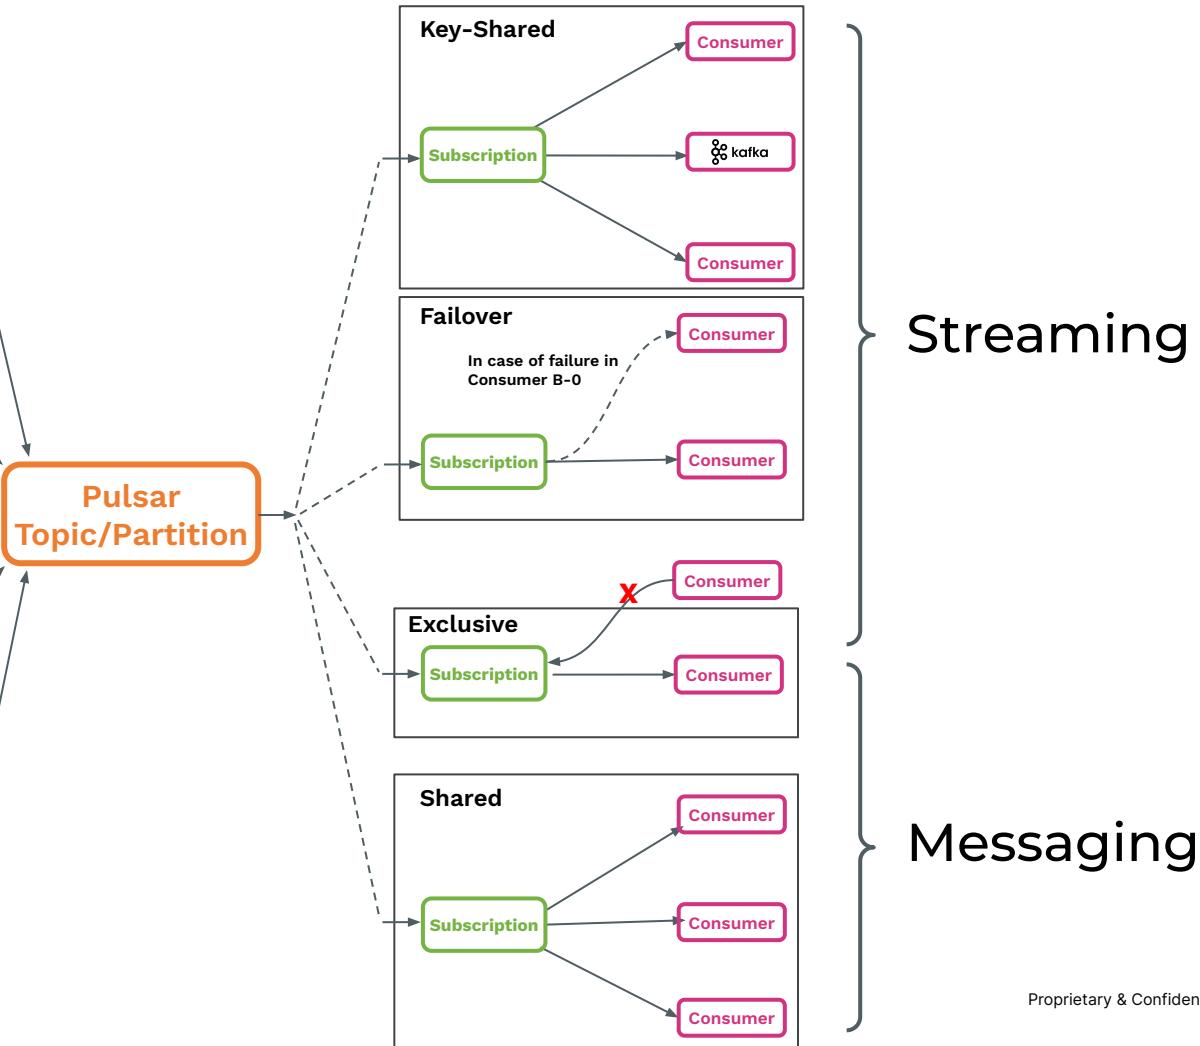
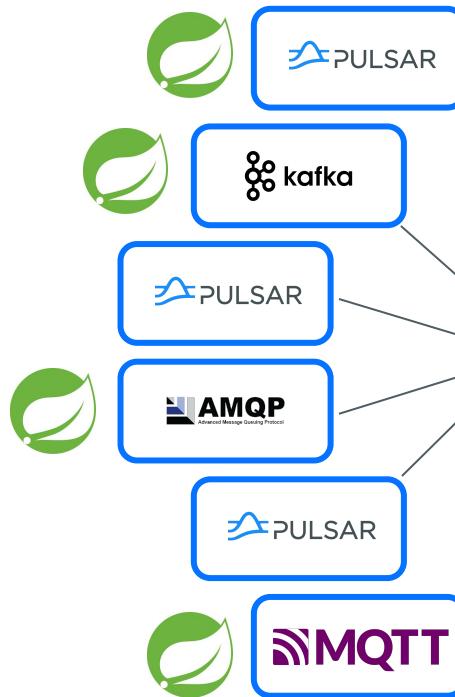
**Guaranteed
Message
Delivery**



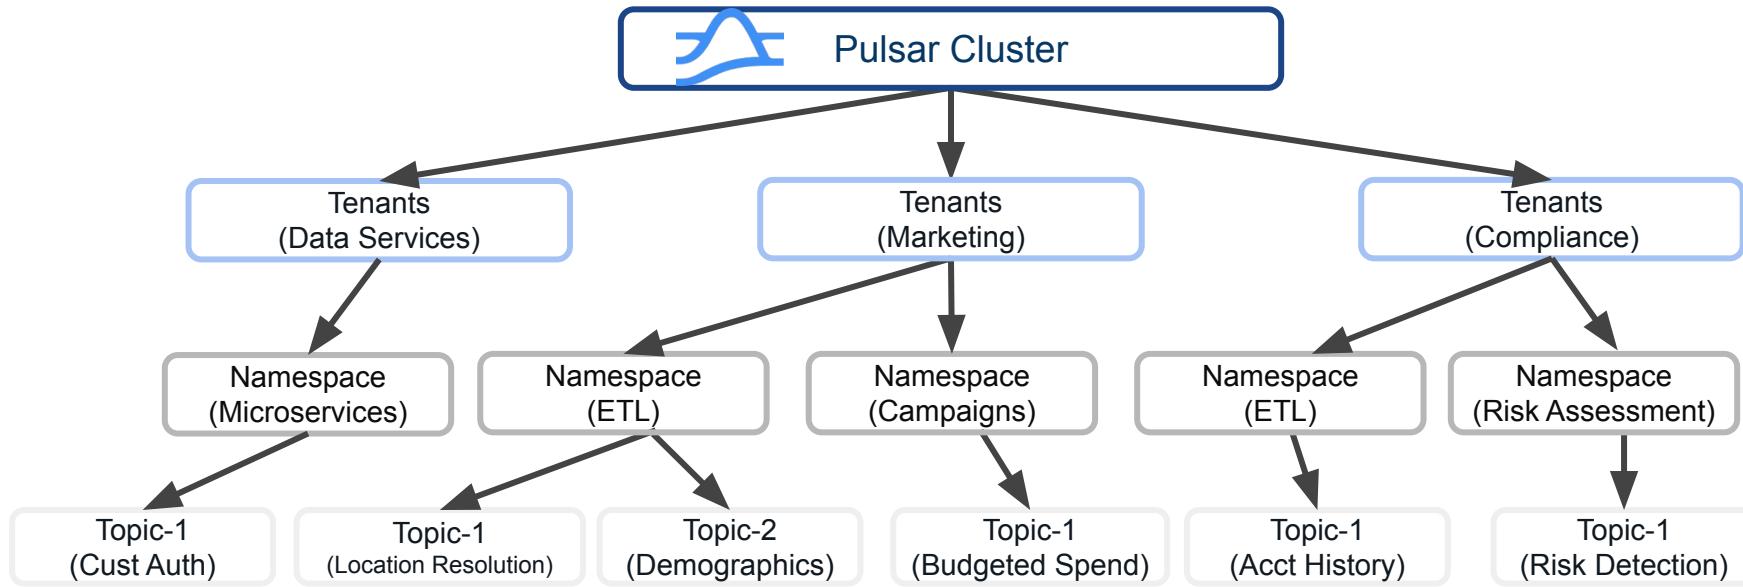
Resiliency



**Infinite
Scalability**



Tenants / Namespaces / Topics

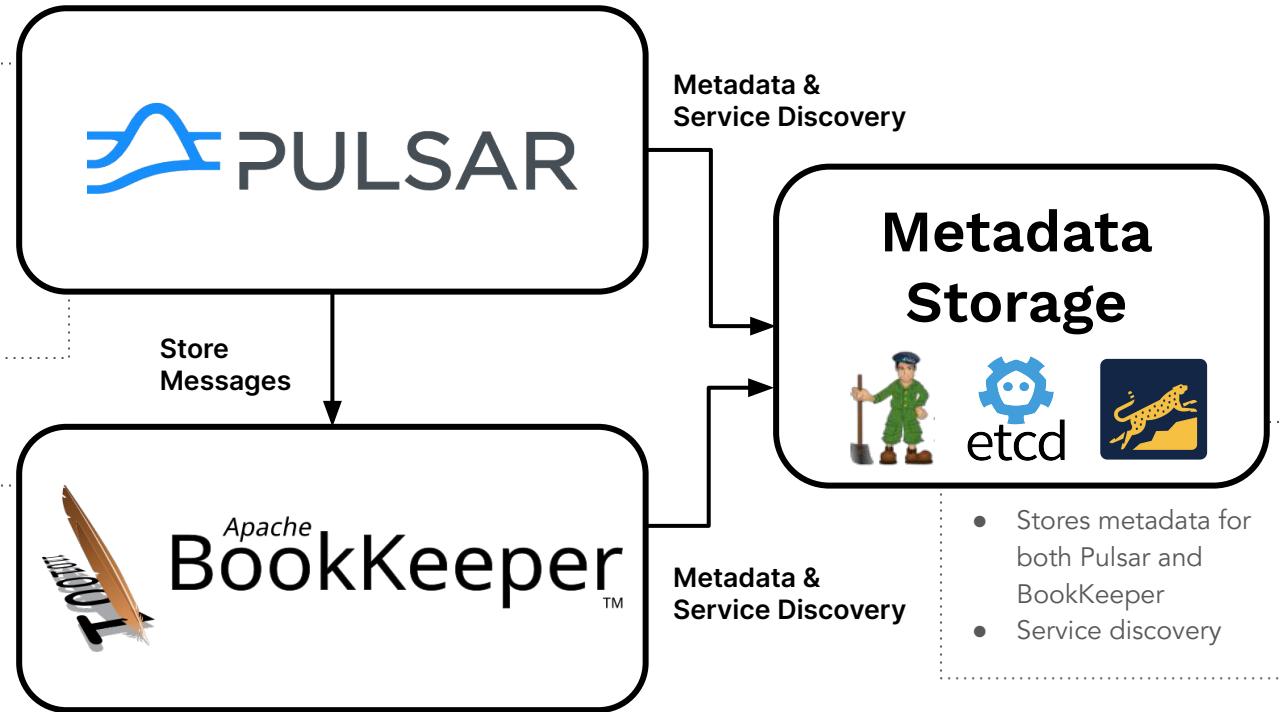


Messages - the basic unit of Pulsar

Component	Description
Value / data payload	The data carried by the message. All Pulsar messages contain raw bytes, although message data can also conform to data schemas.
Key	Messages are optionally tagged with keys, used in partitioning and also is useful for things like topic compaction.
Properties	An optional key/value map of user-defined properties.
Producer name	The name of the producer who produces the message. If you do not specify a producer name, the default name is used.
Sequence ID	Each Pulsar message belongs to an ordered sequence on its topic. The sequence ID of the message is its order in that sequence.

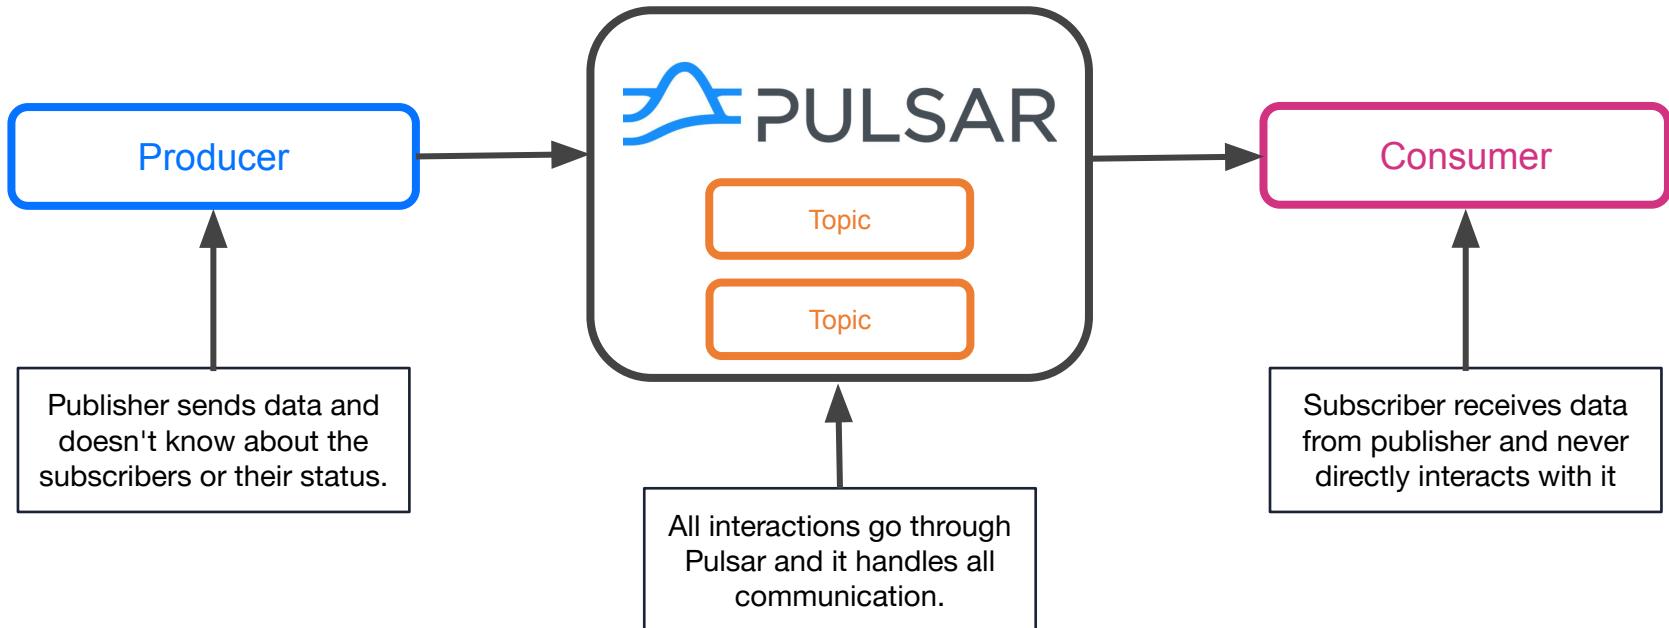
Pulsar Cluster

- “Brokers”
- Handles message routing and connections
- Stateless, but with caches
- Automatic load-balancing
- Topics are composed of multiple segments



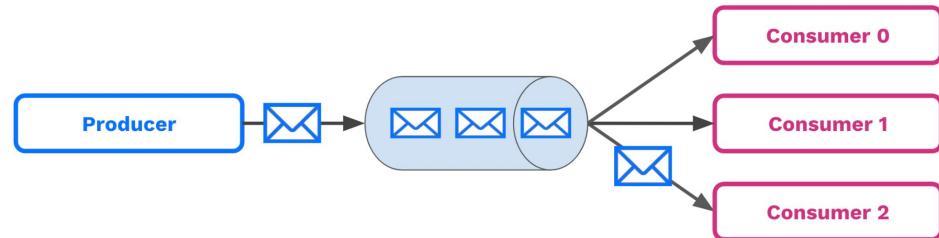
- “Bookies”
- Stores messages and cursors
- Messages are grouped in segments/ledgers
- A group of bookies form an “ensemble” to store a ledger

Producer-Consumer

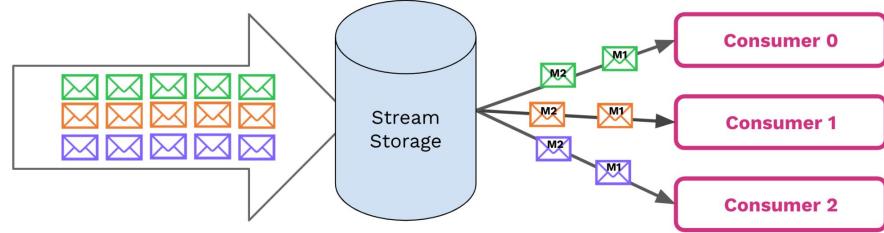


Apache Pulsar: Messaging vs Streaming

Message Queueing - Queueing systems are ideal for work queues that do not require tasks to be performed in a particular order.



Streaming - Streaming works best in situations where the order of messages is important.



Pulsar Subscription Modes

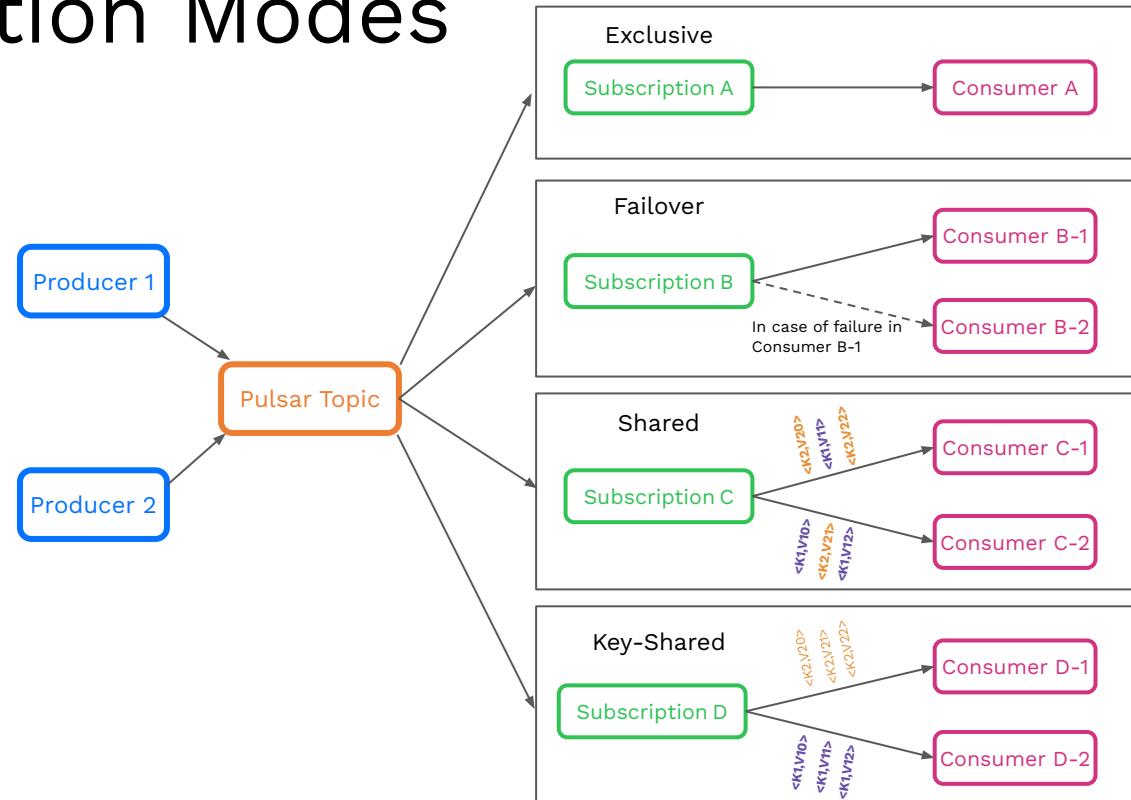
Different subscription modes have different semantics:

Exclusive/Failover -

guaranteed order, single active consumer

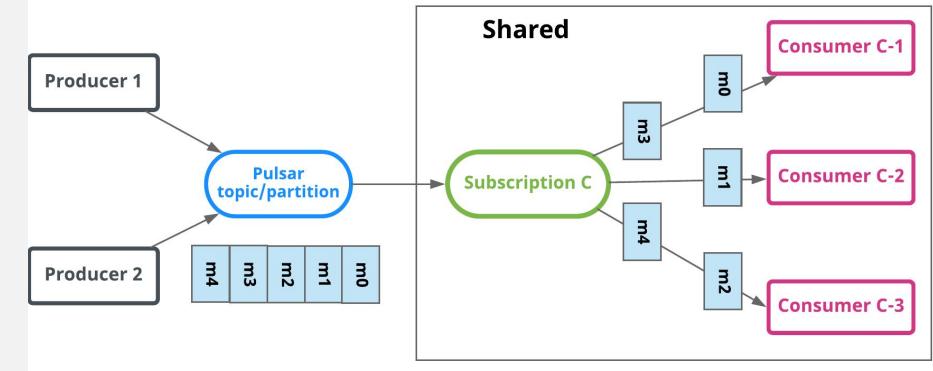
Shared - multiple active consumers, no order

Key_Shared - multiple active consumers, order for given key



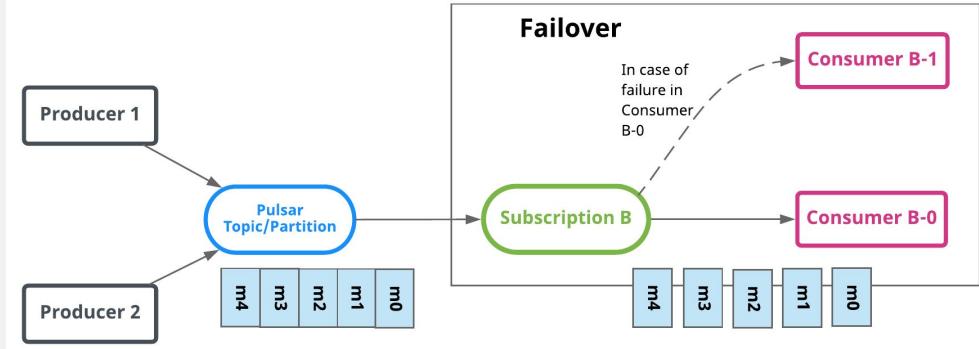
Flexible Pub/Sub API for Pulsar - Shared

```
Consumer consumer = client.newConsumer()  
    .topic("my-topic")  
    .subscriptionName("work-q-1")  
    .subscriptionType(SubType.Shared)  
    .subscribe();
```



Flexible Pub/Sub API for Pulsar - Failover

```
Consumer consumer = client.newConsumer()  
    .topic("my-topic")  
    .subscriptionName("stream-1")  
    .subscriptionType(SubType.Failover)  
    .subscribe();
```



StreamNative Pulsar ecosystem



Protocol Handlers



Client Libraries



Connectors (Sources & Sinks)



Pulsar Functions (Lightweight Stream Processing)



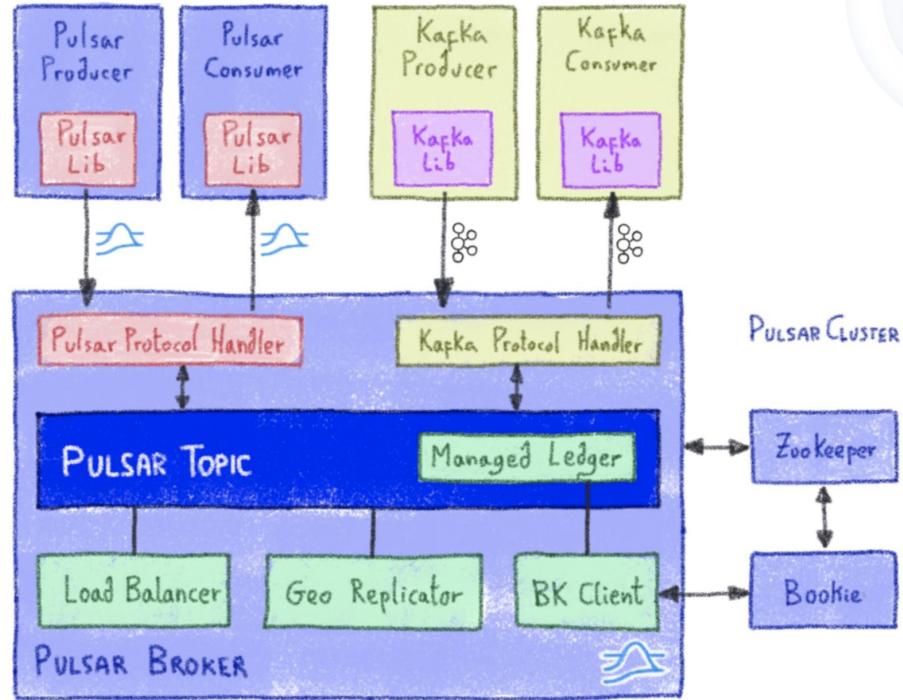
Processing Engines



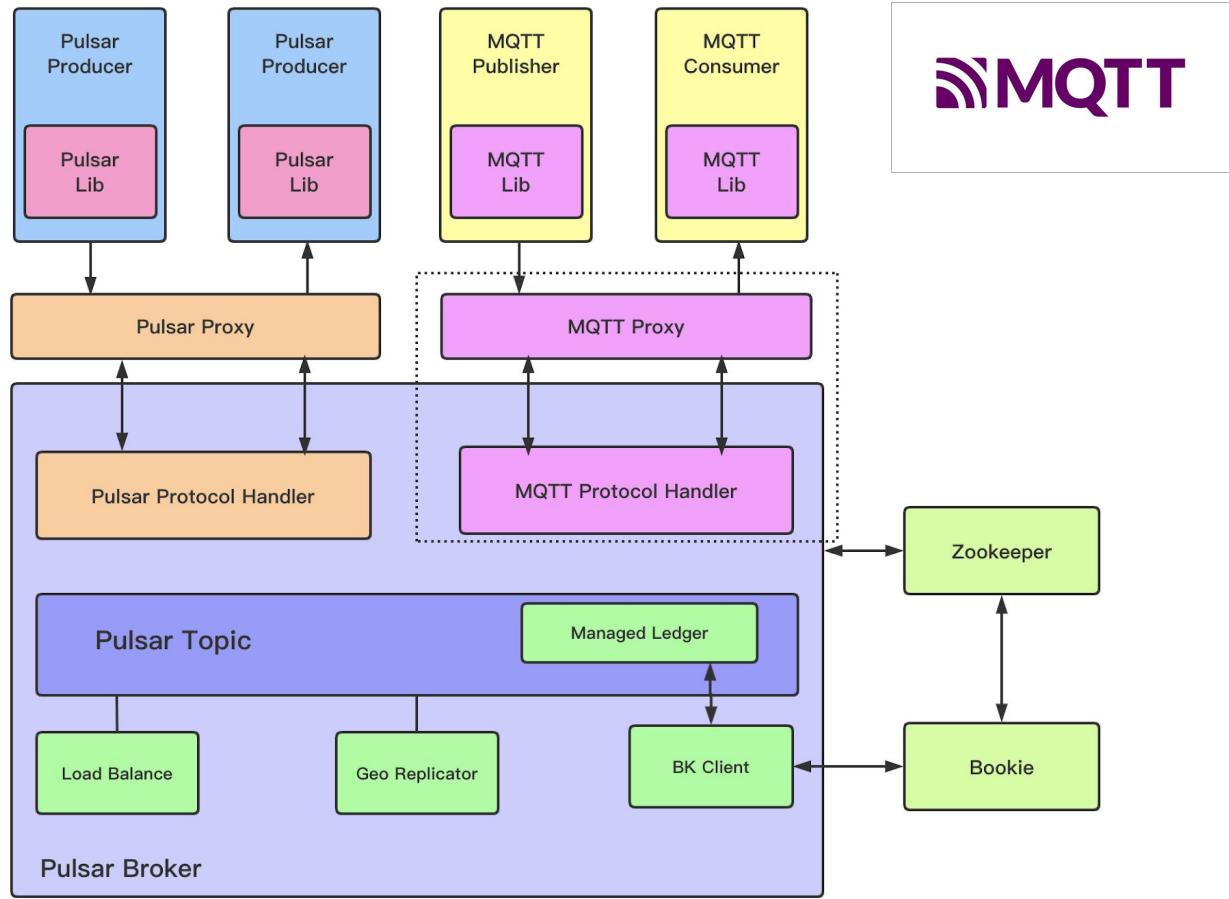
Data Offloaders (Tiered Storage)



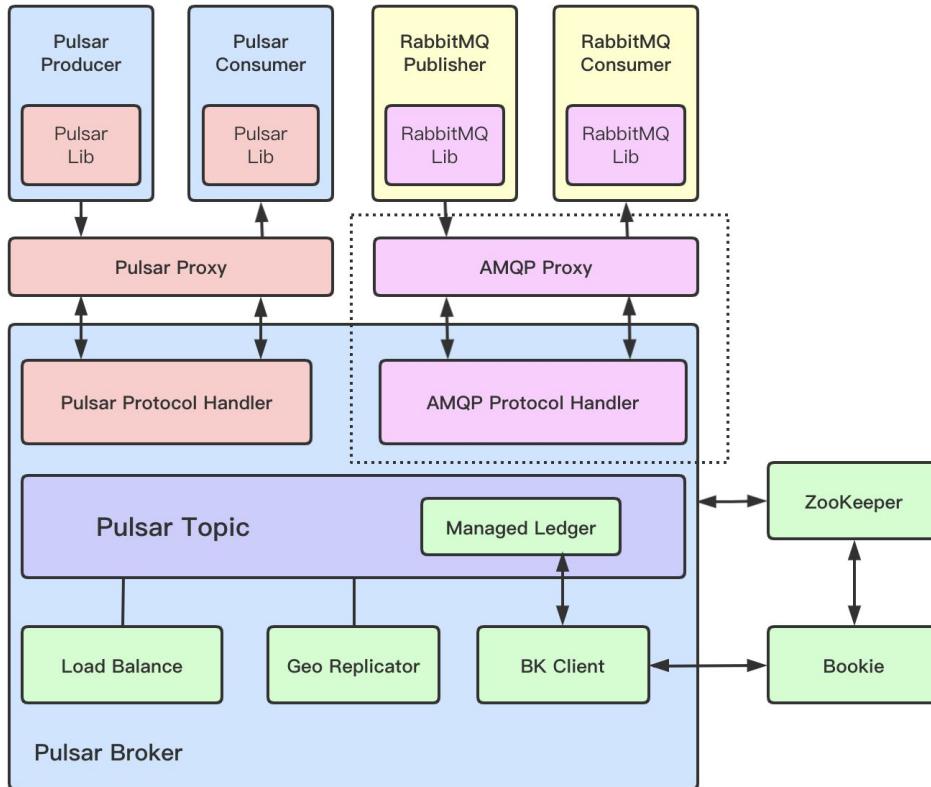
Kafka On Pulsar (KoP)



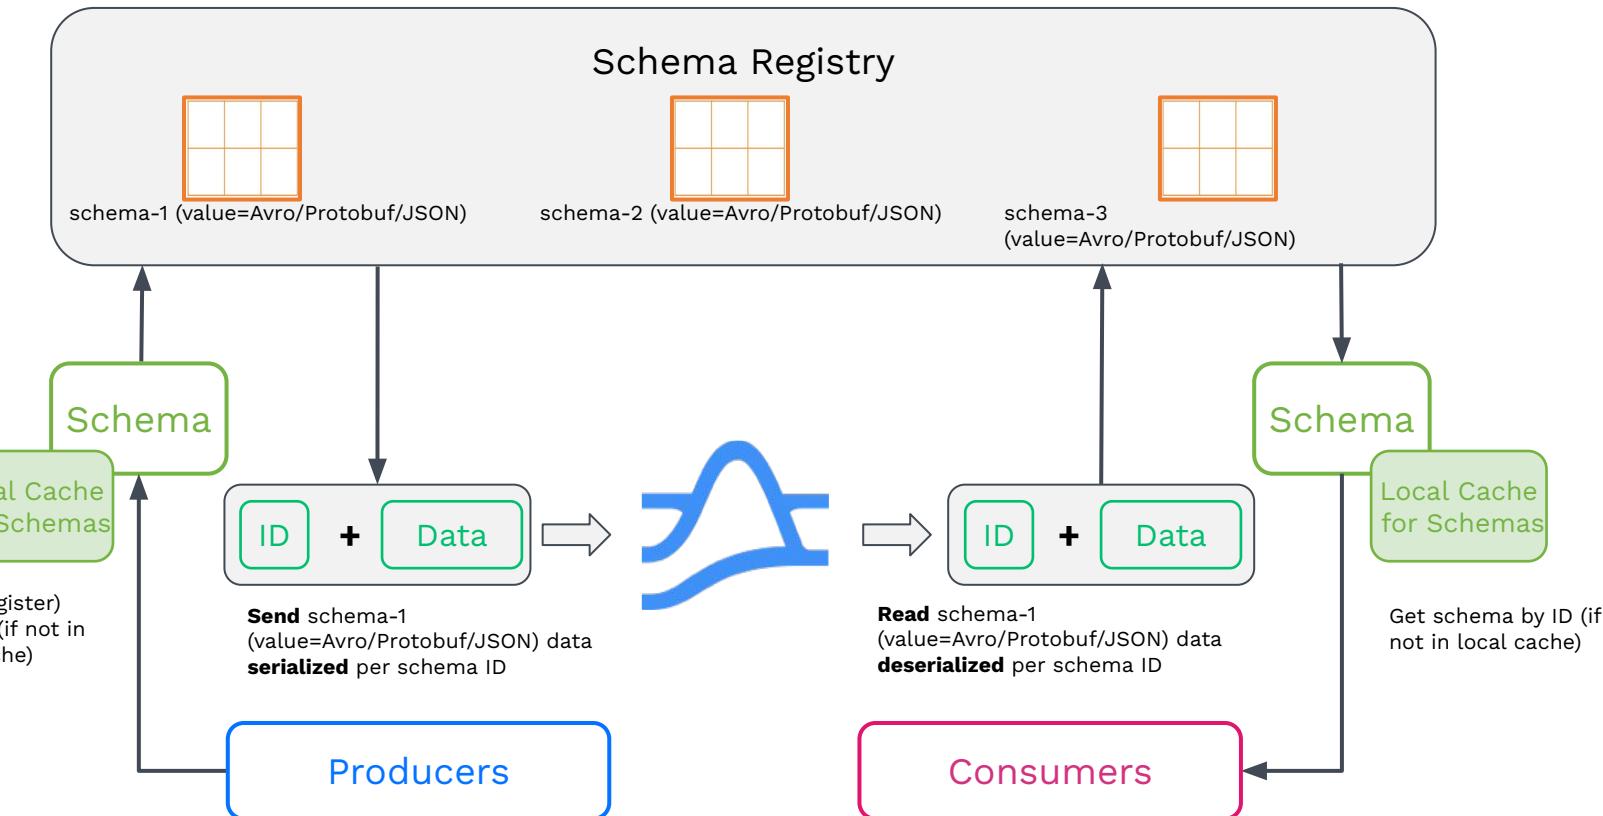
MQTT On Pulsar (MoP)



AMQP On Pulsar (AoP)



Schema Registry



Building Real-Time Requires a Team



Pulsar - Spring



<https://docs.spring.io/spring-pulsar/docs/current/reference/html/>

Pulsar - Spring - Code

```
@Autowired  
private PulsarTemplate<Observation> pulsarTemplate;  
  
this.pulsarTemplate.setSchema(Schema.JSON(Observation.class));  
  
MessageId msgid = pulsarTemplate.newMessage(observation)  
    .withMessageCustomizer(mb -> mb.key(uuidKey.toString()))  
    .send();  
  
@PulsarListener(subscriptionName = "aq-spring-reader", subscriptionType = Shared,  
schemaType = SchemaType.JSON, topics = "persistent://public/default/aq-pm25")  
void echoObservation(Observation message) {  
    this.log.info("PM2.5 Message received: {}", message);  
}
```



Pulsar - Spring - Configuration



```
spring:
  pulsar:
    client:
      service-url: pulsar+ssl://sn-academy.sndevadvocate.snio.cloud:6651
      auth-plugin-class-name: org.apache.pulsar.client.impl.auth.oauth2.AuthenticationOAuth2
      authentication:
        issuer-url: https://auth.streamnative.cloud/
        private-key: file:///scr/sndevadvocate-tspann.json
        audience: urn:sn:pulsar:sndevadvocate:my-instance
    producer:
      batching-enabled: false
      send-timeout-ms: 90000
      producer-name: airqualityjava
      topic-name: persistent://public/default/airquality
```

Spring - Pulsar as Kafka

```
@Bean  
public KafkaTemplate<String, Observation> kafkaTemplate() {  
    KafkaTemplate<String, Observation> kafkaTemplate =  
        new KafkaTemplate<String, Observation>(producerFactory());  
    return kafkaTemplate;  
}
```



```
ProducerRecord<String, Observation> producerRecord = new ProducerRecord<>(topicName,  
    uuidKey.toString(),  
    message);  
  
kafkaTemplate.send(producerRecord);
```



StreamNative

<https://www.baeldung.com/spring-kafka>

Spring - MQTT - Pulsar

```
@Bean
public IMqttClient mqttClient(
    @Value("${mqtt.clientId}") String clientId,
    @Value("${mqtt.hostname}") String hostname,
    @Value("${mqtt.port}") int port)
throws MqttException {
    IMqttClient mqttClient = new MqttClient(
        "tcp://" + hostname + ":" + port, clientId);
    mqttClient.connect(mqttConnectOptions());
    return mqttClient;
}
MqttMessage mqttMessage = new MqttMessage();
mqttMessage.setPayload(DataUtility.serialize(payload));
mqttMessage.setQos(0);
mqttMessage.setRetained(true);
mqttClient.publish(topicName, mqttMessage);
```



Spring - AMQP - Pulsar

@Bean

```
public CachingConnectionFactory  
connectionFactory() {  
    CachingConnectionFactory ccf =  
new CachingConnectionFactory();  
    ccf.setAddresses(serverName);  
    return ccf;  
}
```

```
rabbitTemplate.convertAndSend(topicName,  
DataUtility.serializeToJson(observation));
```



Reactive Spring - Pulsar

2 usages

@Autowired

```
ReactivePulsarTemplate<String> reactivePulsarTemplate;

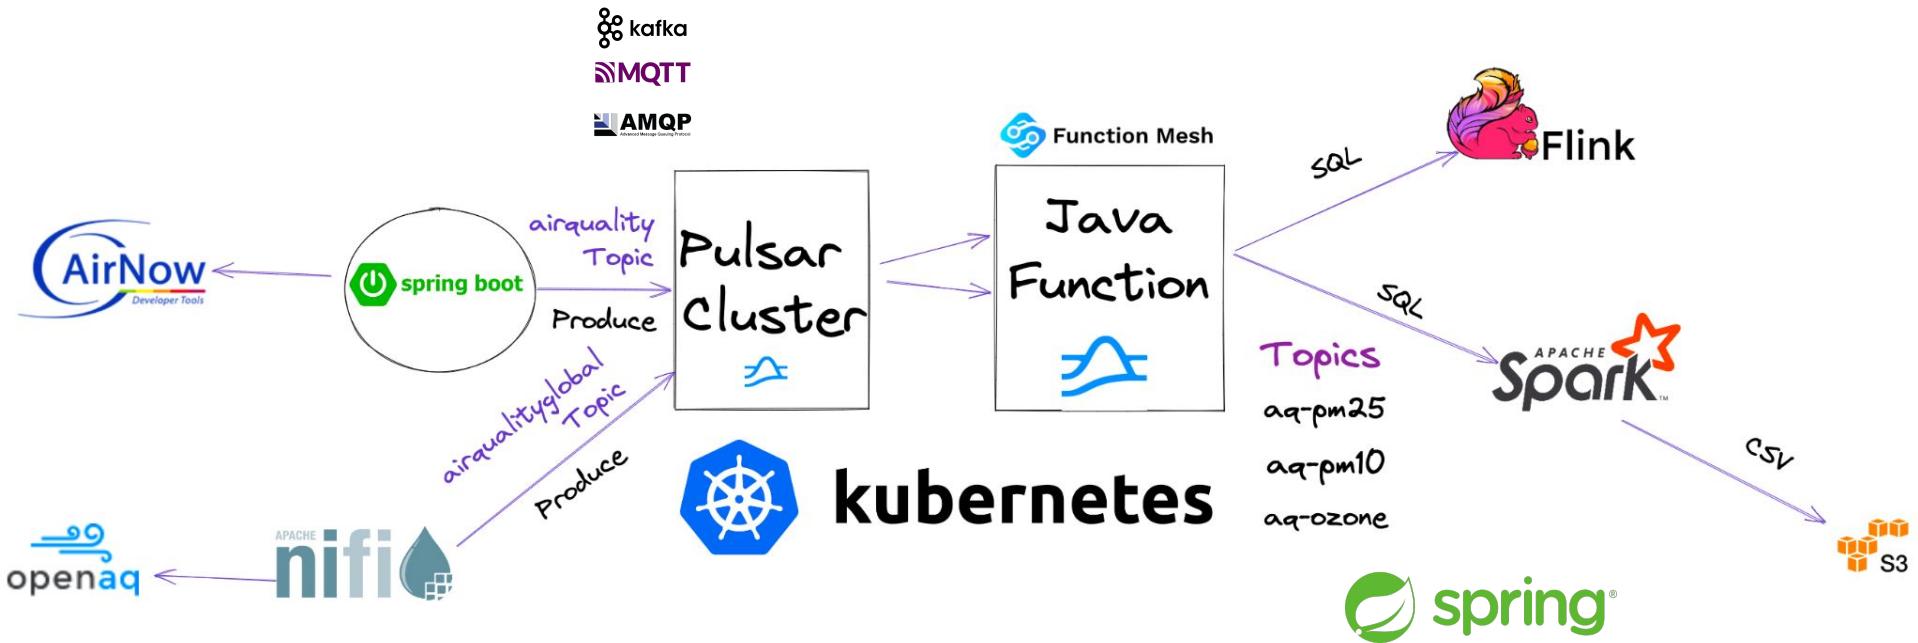
@Timothy Spann *
@scheduled(initialDelay = 1000, fixedRate = 1000)
public void getRows() {
    reactivePulsarTemplate.setSchema(Schema.STRING);

    for (int rowCounter = 0; rowCounter < 20; rowCounter++) {
        System.out.println("sending " + rowCounter);
        reactivePulsarTemplate.newMessage(messageBuilder())
            .withMessageCustomizer((mb) -> mb.key(UUID.randomUUID().toString()))
            .withSenderCustomizer((sc) -> sc.accessMode(ProducerAccessMode.Shared))
            .withSenderCustomizer((sc2) -> sc2.producerName("ReactiveProducer"))
            .withSenderCustomizer((sc3) -> sc3.sendTimeout(Duration.ofSeconds(60L)))
            .withSenderCustomizer((sc4) -> sc4.maxInflight(i: 100))
            .send().subscribe();
    }
}
```

Reactive Spring - Pulsar

```
@ReactivePulsarListener(topics = "persistent://public/default/reactivefaker", stream = true)
Flux<MessageResult<Void>> listen(Flux<Message<String>> messages) {
    return messages
        .doOnNext((msg) -> System.out.println("Stream Received: " + msg.getValue()))
        .map(MessageResult::acknowledge);
}
```

REST + Spring Boot + Pulsar + Friends



FLiP Stack Weekly



<https://bit.ly/32dAJft>



This week in Apache Flink, Apache Pulsar, Apache NiFi, Apache Spark, Java and Open Source friends.

Demo

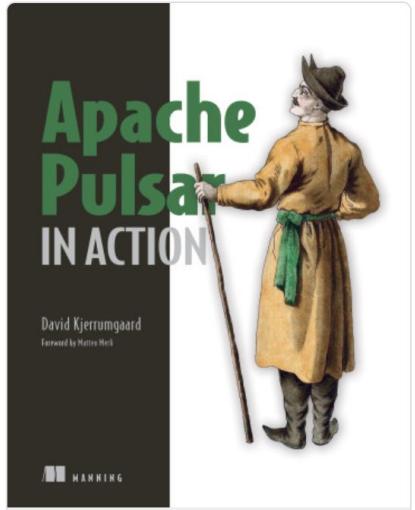


Spring + Pulsar References

- <https://streamnative.io/blog/engineering/2022-11-29-spring-into-pulsar-part-2-spring-based-microservices-for-multiple-protocols-with-apache-pulsar/>
- <https://streamnative.io/blog/release/2022-09-21-announcing-spring-for-apache-pulsar/>
- <https://docs.spring.io/spring-pulsar/docs/current-SNAPSHOT/reference/html/>
- <https://spring.io/blog/2022/08/16/introducing-experimental-spring-support-for-apache-pulsar>
- <https://medium.com/@tspann/using-the-new-spring-boot-apache-pulsar-integration-8a38447dce7b>

Spring Things

- <https://spring.io/guides/gs/spring-boot/>
- <https://spring.io/projects/spring-amqp/>
- <https://spring.io/projects/spring-kafka/>
- <https://github.com/spring-projects/spring-integration-kafka>
- <https://github.com/spring-projects/spring-integration>
- <https://github.com/spring-projects/spring-data-relational>
- <https://github.com/spring-projects/spring-kafka>
- <https://github.com/spring-projects/spring-amqp>



Apache Pulsar in Action

Please enjoy David's complete book which is the ultimate guide to Pulsar.



Tim Spann

Developer Advocate
at StreamNative



<https://www.linkedin.com/in/timothyspann>



@PaaSDev



<https://github.com/tspannhw>

Notices

Apache Pulsar™

Apache®, Apache Pulsar™, Pulsar™, Apache Flink®, Flink®, Apache Spark®, Spark®, Apache NiFi®, NiFi® and the logo are either registered trademarks or trademarks of the [Apache Software Foundation](#) in the United States and/or other countries. No endorsement by The Apache Software Foundation is implied by the use of these marks.

Copyright © 2021-2022 The Apache Software Foundation. All Rights Reserved. Apache, Apache Pulsar and the Apache feather logo are trademarks of The Apache Software Foundation.